


MatrixPolicy step for the reported on-policy configuration


rational coefficients are AdamW-trained and observed; policy actions target RLB matrix blocks

RLB matrix lane: A_l, B_l

sequential substeps; separate optimizer states $A_g \leftarrow e^{-\ell} A_g, B_g \leftarrow e^{-\ell} B_g$



q^R observed
AdamW-only bypass lane



No MatrixPolicy update arrow touches rational coefficients, attention, embeddings, norms, or ordinary weights.