

Rational Matrix Policy Optimization for Rational Feed-Forward Language Models

Anonymous authors
Paper under double-blind review

Abstract

1 Introduction

2 Background and Related Work

Adam-style adaptive methods normalize coordinatewise gradient moments and remain the default optimizer family for many Transformer language-model runs (Kingma & Ba, 2015). Recent optimizer papers have raised the evidence standard for new optimizers: fair tuning over baseline families, throughput and memory accounting, speed-to-target curves, stability reporting, and transfer beyond the selection setting are now expected. MatrixPolicy should be judged against that standard, not against a single default learning rate.

Rational feed-forward layers are a different axis from generic optimizer design. They replace a fixed elementwise nonlinearity with learned rational maps applied to normalized groups. This gives the optimizer access to quantities that do not exist in a SiLU FFN: group radii, rational-coefficient pressures, input/output matrix roles, and a positive gauge relating the adjacent matrices. MatrixPolicy is therefore a structure-specific optimizer. It is complementary to broad optimizer families such as AdamW, Muon-style matrix updates, schedule-free methods, factored adaptive methods, and matrix-preconditioned methods. Our experiments use matched protocols and stability accounting so that the optimizer comparison is about the RLB-specific update rather than an unmatched training recipe.

3 Rational Matrix Policy Optimization

Let $\mathcal{L}(\theta)$ be the language-model training objective. For layer l , the RLB block projects the residual-stream input $x_l \in \mathbb{R}^d$ into G groups of width m ,

$$z_l = A_l x_l, \quad z_{l,g} = \text{group}_g(z_l), \quad (1)$$

normalizes each group by its RMS radius,

$$r_{l,g} = \sqrt{m^{-1} \|z_{l,g}\|_2^2 + \epsilon_{\text{rms}}}, \quad u_{l,g} = z_{l,g} / r_{l,g}, \quad (2)$$

and applies a learned rational map in normalized coordinates,

$$h_{l,g} = r_{l,g} R_{l,g}(u_{l,g}), \quad y_l = B_l \text{concat}_{g=1}^G h_{l,g}. \quad (3)$$

The input matrix A_l selects rational domains, the coefficients of $R_{l,g}$ shape the group response, and the output matrix B_l recombines rational features. MatrixPolicy uses this separation by partitioning parameters as

$$\theta = \theta_{\text{base}} \cup \theta_{\text{coeff}} \cup \{A_l, B_l\}_{l=1}^L. \quad (4)$$

The base parameters use the same optimizer as the non-rational baseline. Rational coefficients use a coefficient rule. MatrixPolicy is applied only to A_l and B_l .

In the homogeneous case $\epsilon_{\text{rms}} = 0$, each group has a positive scale gauge:

$$A_{l,g} \mapsto a_g A_{l,g}, \quad B_{l,g} \mapsto a_g^{-1} B_{l,g}, \quad a_g > 0. \quad (5)$$

The represented function is unchanged because $z_{l,g}$ and $r_{l,g}$ scale together, $u_{l,g}$ is unchanged, and the inverse scale in B_l cancels the scale in $h_{l,g}$. With a nonzero RMS floor, this symmetry is approximate but remains locally controlled under mild nondegeneracy assumptions (Appendix B). This matters because optimizer state is stored in coordinates, not in the quotient space of represented functions.

At step t , each matrix role $r \in \{\text{in}, \text{out}\}$ receives a role-depth factor $\rho_r(l)$ and an on-policy statistic factor $s_{\text{stat}}(l, r, t)$. The matrix update is

$$M_{l,r}^{t+1} = M_{l,r}^t + \Delta_{\text{adapt}}(M_{l,r}; \eta_t a_{\text{mat}}(l, r, t)[1 - \mu(l, r, t)]) + \Delta_{\text{mat}}(M_{l,r}; \eta_t a_{\text{matnorm}}\mu(l, r, t)), \quad (6)$$

where $M_{l,\text{in}} = A_l$, $M_{l,\text{out}} = B_l$, η_t is the shared base learning rate, Δ_{adapt} is an Adam-style adaptive step, and Δ_{mat} is a short early matrix-normalized step. The mixture coefficient $\mu(l, r, t)$ is active only early in training. After the matrix update, MatrixPolicy computes a target log-ratio for each group and applies the bounded rebalance

$$A_{l,g} \leftarrow e^{\ell_g} A_{l,g}, \quad B_{l,g} \leftarrow e^{-\ell_g} B_{l,g}. \quad (7)$$

The update changes the representative used by the optimizer while preserving, or with $\epsilon_{\text{rms}} > 0$ approximately preserving, the represented block function. Full hyperparameter definitions are in Appendix A.

Algorithm 1: Rational Matrix Policy step. Given gradients $\nabla \mathcal{L}(\theta^t)$ and base learning rate η_t :

1. Update exponential moving averages of rational-group pressures.
2. Compute role-depth and group-stat multipliers for each RLB matrix.
3. Update non-rational parameters with the base optimizer and update rational coefficients with the coefficient rule.
4. Update each A_l and B_l with the MatrixPolicy mixture step.
5. Apply the bounded positive-gauge rebalance to each rational group.

4 Experiments

5 Discussion

6 Conclusion

References

Diederik P. Kingma and Jimmy Ba. Adam: A Method for Stochastic Optimization. In International Conference on Learning Representations, 2015. URL <https://mlanthyology.org/iclr/2015/kingma2015iclr-adam/>.

A MatrixPolicy Details

Role-depth factors. For normalized layer depth $d_l = l/(L - 1)$,

$$\rho_{\text{in}}(l) = \text{clip}(1 - 0.50(d_l - 0.5), 0.55, 1.40), \quad (8)$$

$$\rho_{\text{out}}(l) = \text{clip}(1 + 1.00(d_l - 0.5), 0.55, 1.40), \quad (9)$$

$$a_{\text{role}}(l, r) = \max(0.10, 1 + 1.20(\rho_r(l) - 1)). \quad (10)$$

The local adaptive scale is

$$a_{\text{mat}}(l, r, t) = \text{clip}(3.0 a_{\text{role}}(l, r) a_{\text{stat}}(l, r, t), 0.40, 4.0). \quad (11)$$

On-policy group statistics. For group g , MatrixPolicy tracks relative gradient pressures

$$p_{\text{in},g} = \frac{\text{rms}(\nabla A_{l,g})}{\text{rms}(A_{l,g})}, \quad p_{\text{out},g} = \frac{\text{rms}(\nabla B_{l,g})}{\text{rms}(B_{l,g})}, \quad p_{\text{rat},g} = \text{rms}(\nabla \theta_{\text{coeff},g}). \quad (12)$$

With EMAs $q_{\text{in},g}$, $q_{\text{out},g}$, and $q_{\text{rat},g}$, we define

$$\pi_g = \log q_{\text{in},g} - \log q_{\text{out},g}, \quad \alpha_g = \log q_{\text{rat},g} - \frac{1}{2}(\log q_{\text{in},g} + \log q_{\text{out},g}). \quad (13)$$

Group-stat preconditioning. The group preconditioner is centered and clipped:

$$c_g = c_{\text{gain},g} c_{\text{pressure},g} c_{\text{activity},g}, \quad c_g \leftarrow c_g / \text{geomean}(c), \quad c_g \leftarrow \text{clip}(c_g, 0.75, 1.35). \quad (14)$$

The factors are

$$c_{\text{gain},g} = \left(\frac{\text{geomean}(k)}{k_g} \right)^{0.20}, \quad c_{\text{pressure},g} = \exp(0.10 d_{\text{pressure},g}), \quad (15)$$

where $d_{\text{pressure},g} = -\pi_g$ for A_l and $+\pi_g$ for B_l . The activity damping term is

$$c_{\text{activity},g} = \exp[-0.20 \text{ReLU}((\alpha_g - 0.05)/0.45)]. \quad (16)$$

Gauge rebalance. Let

$$\gamma_g = \log \text{rms}(A_{l,g}) - \log \text{rms}(B_{l,g}). \quad (17)$$

The target log-ratio is

$$\tau_g = \frac{1}{2}(\log d_g^{\text{gain}} - \log o_g^{\text{gain}}) + \beta_{\text{pressure}}(\log q_{\text{in},g} - \log q_{\text{out},g}). \quad (18)$$

The applied correction is

$$\ell_g = \frac{1}{2}(\tau_g - \gamma_g), \quad \ell_g \leftarrow \text{clip}(s(t, l) a_g^{\text{activity}} \ell_g, -\ell_{\text{max}}, \ell_{\text{max}}). \quad (19)$$

The transform $A_{l,g} \leftarrow e^{\ell_g} A_{l,g}$ and $B_{l,g} \leftarrow e^{-\ell_g} B_{l,g}$ preserves the homogeneous RLB function exactly and the stabilized implementation up to the RMS-floor perturbation described in Appendix B.

B A Conditional Quotient-Optimization Separation

B.1 Quotient Geometry

For one RLB layer, suppress the layer index and write

$$z_g = A_g x, \quad (20)$$

$$r_g = \sqrt{m^{-1} \|z_g\|_2^2 + \epsilon_{\text{rms}}}, \quad (21)$$

$$u_g = z_g / r_g, \quad (22)$$

$$h_g = r_g R_g(u_g), \quad (23)$$

$$y = B \text{concat}_{g=1}^G h_g. \quad (24)$$

For $a_g > 0$, define

$$T_a(A_g, B_g) = (a_g A_g, a_g^{-1} B_g), \quad g = 1, \dots, G. \quad (25)$$

Equivalently, with $s_g = \log a_g$,

$$T_s(A_g, B_g) = (e^{s_g} A_g, e^{-s_g} B_g). \quad (26)$$

Let $\mathcal{P}$ be the local RLB-matrix parameter space, let $\mathcal{G} = (\mathbb{R}, +)^G$, and let $\mathcal{Q} = \mathcal{P}/\mathcal{G}$.

Lemma B.1 (Exact positive gauge). If $\epsilon_{\text{rms}} = 0$, then $T_s(A, B)$ and (A, B) represent the same layer map for all $s \in \mathbb{R}^G$.

Proof. For one group, $A'_g = e^{s_g} A_g$ gives $z'_g = e^{s_g} z_g$. Since $\epsilon_{\text{rms}} = 0$, $r'_g = e^{s_g} r_g$, hence $u'_g = u_g$ and $h'_g = e^{s_g} h_g$. The corresponding output block satisfies $B'_g h'_g = e^{-s_g} B_g e^{s_g} h_g = B_g h_g$. Summing over groups gives the claim. $\square$

Assumption B.1 (Local quotient chart). In a neighborhood of the trajectory, parameters can be written as

$$\theta = \theta(q, s), \quad q \in \mathbb{R}^d, \quad s \in \mathbb{R}^G, \quad (27)$$

where q is a local coordinate on $\mathcal{Q}$ and s is the log-gauge coordinate. The objective satisfies

$$L(\theta(q, s)) = F(q). \quad (28)$$

For the exact statements, $\epsilon_{\text{rms}} = 0$. Nonzero RMS stabilization is included in the residual model of Assumption B.5.

Assumption B.2 (Local quadratic model). The quotient objective is

$$F(q) = F_\star + \frac{1}{2}(q - q_\star)^\top H(q - q_\star), \quad (29)$$

where $H \succ 0$ and $0 < \mu I \preceq H \preceq LI < \infty$.

Lemma B.2 (Gauge tangent orthogonality). Under Lemma B.1, for any differentiable loss depending on the RLB matrices only through the represented function,

$$\langle \nabla_{A_g} L, A_g \rangle - \langle \nabla_{B_g} L, B_g \rangle = 0. \quad (30)$$

Proof. Differentiate $L(e^{s_g} A_g, e^{-s_g} B_g) = L(A_g, B_g)$ at $s_g = 0$. $\square$

B.2 Canonicalized Quotient Preconditioning

For group g , define

$$\gamma_g = \log \text{rms}(A_g) - \log \text{rms}(B_g). \quad (31)$$

Given a target log-ratio τ_g , the unclipped rebalance is

$$\ell_g = \frac{1}{2}(\tau_g - \gamma_g), \quad A_g^+ = e^{\ell_g} A_g, \quad B_g^+ = e^{-\ell_g} B_g. \quad (32)$$

Lemma B.3 (RMS-ratio canonicalization). The update in equation 32 satisfies

$$\log \text{rms}(A_g^+) - \log \text{rms}(B_g^+) = \tau_g. \quad (33)$$

Proof. By homogeneity of RMS under positive scaling,

$$\log \text{rms}(A_g^+) - \log \text{rms}(B_g^+) = \log \text{rms}(e^{\ell_g} A_g) - \log \text{rms}(e^{-\ell_g} B_g) \quad (34)$$

$$= \gamma_g + 2\ell_g = \tau_g. \quad (35)$$

$\square$

Lemma B.4 (Gauge-dependent raw quotient step). Let $w = ab$, $a = e^s \sqrt{w}$, $b = e^{-s} \sqrt{w}$, and $\ell(w) = \frac{1}{2}(w - y)^2$. One Euclidean gradient-descent step in (a, b) with step size η gives

$$w^+ = w - 2\eta w \cosh(2s)(w - y) + O(\eta^2). \quad (36)$$

Proof. Since $\partial \ell / \partial a = (w - y)b$ and $\partial \ell / \partial b = (w - y)a$,

$$w^+ = (a - \eta(w - y)b)(b - \eta(w - y)a) \quad (37)$$

$$= w - \eta(w - y)(a^2 + b^2) + O(\eta^2). \quad (38)$$

Using $a^2 + b^2 = w(e^{2s} + e^{-2s}) = 2w \cosh(2s)$ gives equation 36. $\square$

Assumption B.3 (Frozen quotient preconditioners). During the local phase, MatrixPolicy and AdamW induce fixed positive definite quotient preconditioners P_{MP} and P_{AdamW} satisfying

$$q_{t+1}^{\text{MP}} = q_t^{\text{MP}} - \eta_{\text{MP}} P_{\text{MP}} \nabla F(q_t^{\text{MP}}), \quad (39)$$

$$q_{t+1}^{\text{AdamW}} = q_t^{\text{AdamW}} - \eta_{\text{AdamW}} P_{\text{AdamW}} \nabla F(q_t^{\text{AdamW}}). \quad (40)$$

Here P_{MP} is the quotient preconditioner after gauge rebalance and P_{AdamW} is the quotient preconditioner induced by AdamW's raw-coordinate adaptive state.

For $P \succ 0$, define

$$K(P) = P^{1/2} H P^{1/2}, \quad \kappa(P; H) = \frac{\lambda_{\max}(K(P))}{\lambda_{\min}(K(P))}. \quad (41)$$

Let

$$\eta^*(P) = \frac{2}{\lambda_{\max}(K(P)) + \lambda_{\min}(K(P))}, \quad \rho(P; H) = \frac{\kappa(P; H) - 1}{\kappa(P; H) + 1}. \quad (42)$$

Assumption B.4 (Strict quotient-conditioning advantage).

$$\kappa_{\text{MP}} := \kappa(P_{\text{MP}}; H) < \kappa(P_{\text{AdamW}}; H) =: \kappa_{\text{AdamW}}. \quad (43)$$

B.3 Contraction and Perturbation Bounds

Lemma B.5 (Preconditioned quadratic contraction). Let F satisfy Assumption B.2. For $P \succ 0$, consider $q_{t+1} = q_t - \eta^*(P) P \nabla F(q_t)$. Then

$$F(q_T) - F_* \leq \rho(P; H)^{2T} (F(q_0) - F_*), \quad (44)$$

and the bound is sharp.

Proof. Let $e_t = q_t - q_*$ and $\tilde{e}_t = P^{-1/2} e_t$. Then

$$\tilde{e}_{t+1} = (I - \eta^*(P) K(P)) \tilde{e}_t, \quad 2(F(q_t) - F_*) = \tilde{e}_t^T K(P) \tilde{e}_t. \quad (45)$$

Diagonalize $K(P) = U \Lambda U^T$ and set $v_t = U^T \tilde{e}_t$. The largest multiplier satisfies

$$\max_i |1 - \eta^*(P) \lambda_i| = \frac{\lambda_{\max}(K(P)) - \lambda_{\min}(K(P))}{\lambda_{\max}(K(P)) + \lambda_{\min}(K(P))} = \rho(P; H). \quad (46)$$

Thus

$$2(F(q_{t+1}) - F_*) = \sum_i \lambda_i v_{t+1,i}^2 \quad (47)$$

$$\leq \rho(P; H)^2 \sum_i \lambda_i v_{t,i}^2 \quad (48)$$

$$= 2\rho(P; H)^2 (F(q_t) - F_*). \quad (49)$$

Iteration gives equation 44. Equality holds for an eigenvector of $K(P)$ associated with $\lambda_{\min}(K(P))$ or $\lambda_{\max}(K(P))$. $\square$

Theorem B.1 (Conditional optimization separation). Suppose Assumptions B.1–B.4 hold, and set $\eta_{\text{MP}} = \eta^*(P_{\text{MP}})$ and $\eta_{\text{AdamW}} = \eta^*(P_{\text{AdamW}})$. Then, for every integer $T \geq 1$,

$$\sup_{q_0 \neq q_*} \frac{F(q_T^{\text{MP}}) - F_*}{F(q_0) - F_*} \leq \left(\frac{\kappa_{\text{MP}} - 1}{\kappa_{\text{MP}} + 1} \right)^{2T} < \left(\frac{\kappa_{\text{AdamW}} - 1}{\kappa_{\text{AdamW}} + 1} \right)^{2T} = \sup_{q_0 \neq q_*} \frac{F(q_T^{\text{AdamW}}) - F_*}{F(q_0) - F_*}. \quad (50)$$

Proof. The map $\kappa \mapsto (\kappa - 1)/(\kappa + 1)$ is strictly increasing for $\kappa > 1$. Assumption B.4 gives $\rho(P_{\text{MP}}; H) < \rho(P_{\text{AdamW}}; H)$. Lemma B.5 applied to P_{MP} and P_{AdamW} gives the two displayed bounds. The AdamW equality follows from sharpness in Lemma B.5. $\square$

Assumption B.5 (Uniform local perturbation bound). For a fixed horizon T , the implemented quotient iterates satisfy

$$q_{t+1}^{\text{MP,impl}} = q_t^{\text{MP,impl}} - \eta_{\text{MP}} P_{\text{MP}} \nabla F(q_t^{\text{MP,impl}}) + r_t^{\text{MP}}, \quad (51)$$

$$q_{t+1}^{\text{AdamW,impl}} = q_t^{\text{AdamW,impl}} - \eta_{\text{AdamW}} P_{\text{AdamW}} \nabla F(q_t^{\text{AdamW,impl}}) + r_t^{\text{AdamW}}, \quad (52)$$

with $\|r_t^{\text{MP}}\|_2 + \|r_t^{\text{AdamW}}\|_2 \leq \delta$ for all $0 \leq t < T$.

Proposition B.1 (Stability of a strict finite-time gap). Fix $T \geq 1$ and suppose the idealized dynamics satisfy

$$m_T := F(q_T^{\text{AdamW}}) - F(q_T^{\text{MP}}) > 0. \quad (53)$$

Assume the ideal trajectories remain in a compact region on which F and the ideal one-step maps are Lipschitz. Then there exists $\delta_0 > 0$ such that, whenever Assumption B.5 holds with $\delta < \delta_0$,

$$F(q_T^{\text{AdamW,impl}}) - F(q_T^{\text{MP,impl}}) > \frac{m_T}{2} > 0. \quad (54)$$

Proof. Let Ψ_{MP} and Ψ_{AdamW} be the ideal one-step maps. On a compact neighborhood of the ideal trajectories, let C_Ψ be a common Lipschitz constant. By induction,

$$\|q_t^{\text{MP,impl}} - q_t^{\text{MP}}\| + \|q_t^{\text{AdamW,impl}} - q_t^{\text{AdamW}}\| \leq C_T \delta \quad (55)$$

for a constant C_T depending only on T and C_Ψ . If C_F is a Lipschitz constant for F on the same compact region, then

$$\left| F(q_T^{\text{MP,impl}}) - F(q_T^{\text{MP}}) \right| + \left| F(q_T^{\text{AdamW,impl}}) - F(q_T^{\text{AdamW}}) \right| \leq C_F C_T \delta. \quad (56)$$

Taking $\delta_0 = m_T / (2C_F C_T)$ gives the claim. $\square$