

MATRIXPOLICY FOR ROLE-AWARE OPTIMIZATION OF RATIONAL TRANSFORMER NONLINEARITIES

Anonymous authors

Paper under double-blind review

ABSTRACT

A useful LLM pretraining recipe moves the loss-time frontier by reaching the same validation loss with fewer tokens or less wall-clock time. We target this objective through the Transformer nonlinear sublayer, treating it as an optimizer-visible interface. MatrixPolicy with Rational Latent Basis (RLB) replaces a fixed scalar response with grouped rational responses that expose local gain, coefficient-activity, and matrix-balance statistics to a role-aware update. MatrixPolicy restricts the signal-conditioned update to the sublayer input and output matrices; rational coefficients and the rest of the Transformer remain on AdamW. Across DCLM, FineWeb-Edu, FineWeb, Dolma-sample, and C4 with a fixed 12-layer, width-768 decoder, MatrixPolicy reaches matched validation-loss targets with fewer tokens and lower measured arrival time than SiLU controls under the same protocol. It saves 19.5–22.8% target-arrival tokens versus SiLU with AdamW and 29.2–34.0% versus SiLU with Muon at 100M tokens, and 22.3–26.4% and 6.7–9.8%, respectively, at 300M tokens. [todo: Insert the larger-scale training result here.]

1 INTRODUCTION

Recent LLM optimizer papers frame practical efficiency through loss-time and compute-time frontiers, asking which training recipe reaches a target loss with fewer tokens, FLOPs, or wall-clock time (Shah et al., 2025; Liu et al., 2025; Semenov et al., 2025). This criterion also exposes an activation-design question. In a modern decoder block, the nonlinear sublayer is a learned matrix–nonlinearity–matrix map. The activation sets local gain and curvature, while the adjacent matrices choose and recombine the coordinates on which that response acts. Standard training pipelines leave this local structure implicit to the optimizer. Activation coefficients and adjacent matrices are updated as ordinary parameters, leaving gain, coefficient-activity, and role information inside the sublayer unused by the matrix update.

We make this interface concrete with Rational Latent Basis (RLB) and MatrixPolicy. RLB maps residual states into grouped latent coordinates, normalizes each group by its RMS radius, and applies a trainable real-pole-free P5/Q4 rational response. The resulting response-gain, coefficient-activity, matrix-pressure, and pair-balance signals are detached from the forward graph and passed to MatrixPolicy. MatrixPolicy restricts their influence to the RLB input and output matrices; attention, embeddings, normalization parameters, rational coefficients, and ordinary Transformer tensors remain on AdamW. This local division exposes the quantities created by the nonlinear sublayer while preserving AdamW for the rest of the model.

We evaluate target arrival by measuring the tokens and wall-clock time required to reach a fixed validation-loss target under the same model, corpus, and protocol. In the completed study, the same 12-layer, width-768 causal Transformer is trained on DCLM, FineWeb-Edu, FineWeb, Dolma-sample, and C4. The experiments include RLB-only controls and matched SiLU with AdamW and SiLU with Muon baselines under the same model, data, and measurement protocol. MatrixPolicy reaches shared validation-loss targets earlier at both 100M and 300M token budgets. [todo: Insert the larger-scale training result here.] The contribution is an activation design that exposes optimizer-visible matrix roles, a localized role-aware optimizer for the exposed matrix pair, and matched evidence that this interface reduces target-arrival tokens and wall-clock time across five pretraining corpora.

2 RELATED WORK

Activation design for Transformer nonlinear sublayers has moved from fixed pointwise curves to gated matrix sublayers. Early Transformer-style sublayers used fixed pointwise responses such as GELU or Swish/SiLU (Vaswani et al., 2017; Hendrycks & Gimpel, 2016; Ramachandran et al., 2017). GLU introduced multiplicative feature selection for language modeling (Dauphin et al., 2017), GLU variants such as SwiGLU improved Transformer sublayers in compute-matched settings (Shazeer, 2020), and large decoder models adopted SwiGLU together with pre-normalized Transformer components (Chowdhery et al., 2022; Touvron et al., 2023). These works identify the nonlinear block as the natural comparison point for RLB because the response is coupled to the matrices that prepare and mix it. RLB keeps that sublayer location while replacing the fixed curve with grouped trainable rational responses.

Choosing a trainable rational response builds on earlier evidence that low-degree quotients can serve as learnable nonlinearities. Padé Activation Units learned such quotients end to end as activation functions (Molina et al., 2020), rational neural networks gave approximation efficiency results for smooth-function classes (Boullé et al., 2020), and RAFT evaluated rational activation functions inside Transformer pretraining and fine-tuning (Fang et al., 2023). Recent theory strengthens the motivation by proving expressivity advantages for trainable low-degree rational activations over a broad set of modern fixed activations, including gated and Transformer-style nonlinearities (Tang & Townsend, 2026). RLB therefore follows the rational-activation lineage; MatrixPolicy tests the additional effect of making the surrounding matrices optimizer-visible.

Making the surrounding matrices optimizer-visible connects RLB to a fast-moving LLM optimizer literature. Adam and AdamW remain central baselines (Kingma & Ba, 2015; Loshchilov & Hutter, 2019), while state-efficient or diagonal adaptive methods such as Adafactor and CAME reduce optimizer memory or alter adaptive-state estimates (Shazeer & Stern, 2018; Luo et al., 2023). Other recent methods change the update rule itself. Lion, Schedule-Free methods, and cautious optimizers alter momentum, sign, or schedule behavior (Chen et al., 2023; Defazio et al., 2024; Liang et al., 2026), while Sophia, GaLore, and Adam-mini change curvature, low-rank, or per-block learning-rate structure (Liu et al., 2024; Zhao et al., 2024; Zhang et al., 2025). Practical-efficiency studies evaluate optimizers by the tokens, compute, or time needed to reach matched losses (Shah et al., 2025; Schlotthauer et al., 2025; Semenov et al., 2025; Wen et al., 2026). The matrix restriction in MatrixPolicy connects most directly to matrix-aware optimization. Shampoo uses tensor preconditioners (Gupta et al., 2018), SOAP stabilizes Shampoo-style preconditioning by running Adam in the preconditioner eigenbasis (Vyas et al., 2025), and Muon uses orthogonalized matrix updates for LLM training (Liu et al., 2025). Recent optimizer comparisons emphasize fair tuning, model-scale dependence, and setting-dependent tradeoffs (Zhao et al., 2025; Wen et al., 2026). MatrixPolicy follows this matrix-aware direction, with updates restricted to the RLB input/output matrix pair and conditioned on activation-derived sublayer statistics.

3 METHOD

RLB and MatrixPolicy are designed as a single interface. RLB defines the nonlinear sublayer and the detached summaries it exposes; MatrixPolicy consumes those summaries to update the two adjacent matrices. Rational coefficients and ordinary Transformer parameters remain on the AdamW path in the evaluated configuration.

Throughout the method we use

$$\text{clip}(x, a, b) = \min\{\max\{x, a\}, b\}, \quad (1)$$

$$\text{smoothstep}(a, b, x) = 3\omega^2 - 2\omega^3, \quad \omega = \text{clip}\left(\frac{x-a}{b-a}, 0, 1\right). \quad (2)$$




Figure 1: Activation-optimizer interface for RLB. In layer l , A_l maps the sublayer input into normalized rational groups and B_l recombines the rational features $h_l = \text{concat}_{g=1}^G h_{l,g}$. Forward RLB statistics, gradient pressure/activity EMAs, and pair-balance measurements cross a stop-gradient boundary as policy inputs. MatrixPolicy reads these summaries and writes ΔA_l and ΔB_l ; rational coefficients are AdamW-trained and policy-observed.

3.1 RATIONAL LATENT BASIS

For layer $l \in \{1, \dots, L\}$, let $x_l \in \mathbb{R}^d$ be the input to the nonlinear sublayer. RLB uses an input matrix $A_l \in \mathbb{R}^{H \times d}$ and an output matrix $B_l \in \mathbb{R}^{d \times H}$ with latent width $H = Gm$. Let $A_{l,g} \in \mathbb{R}^{m \times d}$ denote the row block of A_l , and let $B_{l,g} \in \mathbb{R}^{d \times m}$ denote the matching column block of B_l . In the implementation, A_l and B_l are the bias-free `in_proj.weight` and `out_proj.weight` of the nonlinear sublayer. The evaluated runs use an RMS floor $\epsilon_{\text{rms}} = 10^{-6}$. Each group has a trainable P5/Q4 rational response

$$P_{l,g}(u) = \sum_{i=0}^5 a_{l,g,i} u^i, \quad (3)$$

$$Q_{l,g}(u) = 1 + |b_{l,g,1}| |u| + |b_{l,g,2}| u^2 + |b_{l,g,3}| |u|^3 + |b_{l,g,4}| u^4, \quad (4)$$

$$R_{l,g}(u) = P_{l,g}(u)/Q_{l,g}(u). \quad (5)$$

Define the preactivation, RMS radius, normalized coordinate, and grouped rational feature by

$$z_{l,g}(x_l) = A_{l,g} x_l, \quad \rho_{l,g}(x_l) = \sqrt{m^{-1} \|z_{l,g}(x_l)\|_2^2 + \epsilon_{\text{rms}}}, \quad (6)$$

$$u_{l,g}(x_l) = \frac{z_{l,g}(x_l)}{\rho_{l,g}(x_l)}, \quad h_{l,g}(x_l) = \rho_{l,g}(x_l) R_{l,g}(u_{l,g}(x_l)), \quad h_l = \text{concat}_{g=1}^G h_{l,g}. \quad (7)$$

The RLB forward map is the single block expression

$$y_l = B_l h_l = \sum_{g=1}^G B_{l,g} \rho_{l,g}(x_l) R_{l,g} \left(\frac{A_{l,g} x_l}{\rho_{l,g}(x_l)} \right). \quad (8)$$

The scalar functions $P_{l,g}$, $Q_{l,g}$, and $R_{l,g}$ are applied coordinatewise when their input is a vector. Since $Q_{l,g}(u) \geq 1$ for all real u , the implemented response is pole-free on $\mathbb{R}$. The response is piecewise rational because of the absolute-value terms; derivative statistics use the autodiff convention $\text{sign}(0) = 0$ and equations involving $R'_{l,g}$ are interpreted almost everywhere. Numerator and denominator coefficients are initialized by a SiLU fit on $[-5, 5]$ and are then trained by AdamW. The evaluated model uses these groupwise P5/Q4 rational responses.

The forward map in equation 8 separates the normalized coordinate at which the rational response is evaluated from the group radius returned to the residual stream. This radius-factor form exposes

the positive homogeneity behind the matrix-pair coordinate formalized in Appendix B. The floor $\epsilon_{\text{rms}} > 0$ is an additive squared-RMS stabilizer. For fixed curves $R_l = \{R_{l,g}\}_{g=1}^G$, write this block map as $f_l(x; A_l, B_l, R_l)$. For any fixed input x such that $A_{l,g}x \neq 0$ for every group, if $\epsilon_{\text{rms}} = 0$, then for any $\lambda \in \mathbb{R}_{>0}^G$ and $D_l(\lambda) = \text{diag}(\lambda_1 I_m, \dots, \lambda_G I_m)$,

$$f_l(x; A_l, B_l, R_l) = f_l(x; D_l(\lambda)A_l, B_l D_l(\lambda)^{-1}, R_l). \quad (9)$$

This positive matrix-pair rescaling is exact in the unfloored map. The RMS floor perturbs the orbit, and MatrixPolicy follows the resulting approximate coordinate with bounded steps.

3.2 MATRIXPOLICY

MatrixPolicy is the optimizer-side policy attached to the RLB matrix pair (A_l, B_l) . We use the name for the matrix-specific controls in the evaluated optimizer configuration. Other Transformer tensors follow their assigned backbone optimizer, and the rational coefficients contribute activity signals while retaining their AdamW no-decay updates in the reported runs. The policy reads detached summaries. These summaries are produced by the RLB forward hook, by post-global-clipping gradients observed before MatrixPolicy rescales them, and by fixed-grid probes of the rational response. They feed, in order, group-gradient gating, a role/depth matrix step, and reciprocal pair rescaling.

This ordering matters. Let t be the one-based optimizer step after the wrapper increments its counter, T the planned number of steps, and $\xi_t = t/T$. The surrounding optimizer wrapper first constructs detached gradient memories from the current step; the inner matrix optimizer uses those memories to rescale matrix gradients and apply AdamW/Muon updates; the wrapper then measures the updated matrix pair and applies the reciprocal rescale. No summary tensor enters the forward graph. For notation, set $M_{l,\text{in}} = A_l$ and $M_{l,\text{out}} = B_l$, with $\sigma \in \{\text{in}, \text{out}\}$ denoting the two roles. For any tensor V , define

$$\text{rms}(V) = \sqrt{\text{mean}(V^2)}, \quad \text{rms}_\epsilon(V) = \sqrt{\text{mean}(V^2) + \epsilon}. \quad (10)$$

The evaluated implementation uses $\epsilon_{\text{live}} = 10^{-6}$ for cached forward gains and $\epsilon_p = \epsilon_{\text{probe}} = 10^{-8}$ for pressure, matrix-ratio, and fixed-grid probe floors.

The first family of summaries is the on-policy pressure and activity memory. Here “on-policy” means that the wrapper records statistics from the current training trajectory. All gradients in this paragraph are optimizer-visible loss gradients after global clipping and before MatrixPolicy group-gradient scaling. Let $\Theta_{l,g}^R$ collect the trainable rational-response coefficient blocks in group (l, g) ; for the reported P5/Q4 response these are the numerator and denominator coefficient tensors. The count $|\Theta_{l,g}^R|$ is a block count, and each mean is over scalar entries in one block. The instantaneous detached measurements are

$$p_{l,g}^{\text{in}} = \frac{\text{rms}_{\epsilon_p}(\nabla A_{l,g})}{\text{rms}_{\epsilon_p}(A_{l,g})}, \quad p_{l,g}^{\text{out}} = \frac{\text{rms}_{\epsilon_p}(\nabla B_{l,g})}{\text{rms}_{\epsilon_p}(B_{l,g})}, \quad (11)$$

$$p_{l,g}^R = \left[\frac{1}{|\Theta_{l,g}^R|} \sum_{\theta \in \Theta_{l,g}^R} \text{mean}((\nabla \theta)^2) + \epsilon_p \right]^{1/2}. \quad (12)$$

The matrix quantities are relative RMS-gradient proxies, while $p_{l,g}^R$ is raw coefficient-gradient energy. The wrapper stores detached EMA memories q^{in} , q^{out} , and q^R ,

$$q_{l,g,t}^r = 0.95 q_{l,g,t-1}^r + 0.05 p_{l,g,t}^r, \quad r \in \{\text{in}, \text{out}, R\}, \quad (13)$$

initialized from the first observed values. We reserve q in this subsection for these detached policy memories; the rational denominator remains denoted by $Q_{l,g}$. The two contrasts used by all three controls are

$$\pi_{l,g} = \log q_{l,g}^{\text{in}} - \log q_{l,g}^{\text{out}}, \quad \alpha_{l,g} = \log q_{l,g}^R - \frac{1}{2} (\log q_{l,g}^{\text{in}} + \log q_{l,g}^{\text{out}}). \quad (14)$$

The pressure contrast $\pi_{l,g}$ compares the two adjacent matrix roles, and the activity contrast $\alpha_{l,g}$ compares rational-coefficient gradient energy with the geometric mean of those role memories. Appendix C.1 derives the local gradient geometry behind these proxies. We drop the step subscript when the time is clear.

The second family of summaries comes from the RLB forward-statistics hook. The hook samples normalized coordinates and caches per-group response and derivative RMS values,

$$\hat{d}_{l,g}^{\text{live}} = \sqrt{\text{mean}(R'_{l,g}(u_{l,g})^2) + \epsilon_{\text{live}}}, \quad \hat{o}_{l,g}^{\text{live}} = \sqrt{\text{mean}(R_{l,g}(u_{l,g})^2) + \epsilon_{\text{live}}}. \quad (15)$$

The mean is over sampled scalar normalized coordinates for group (l, g) across the most recent refreshed sample, sequence positions, and the m group coordinates. The derivative gain is the input-role gain proxy, and the response gain is the output-role gain proxy. If a cache is unavailable at a step, the implementation sets the corresponding gain multiplier to identity while keeping the pressure/activity memories active.

The third family supplies the post-step pair rescale. Let $\mathcal{U} = \{u_k\}_{k=1}^{129}$ be a uniform probe grid over $[-5, 5]$, and define $\text{rms}_{\mathcal{U}}(\varphi) = \sqrt{|\mathcal{U}|^{-1} \sum_{u_k \in \mathcal{U}} \varphi(u_k)^2}$. The fixed-grid probe producer refreshes every ten optimizer steps and emits

$$\hat{d}_{l,g}^{\text{probe}} = \sqrt{\text{rms}_{\mathcal{U}}(R'_{l,g})^2 + \epsilon_{\text{probe}}}, \quad \hat{o}_{l,g}^{\text{probe}} = \sqrt{\text{rms}_{\mathcal{U}}(R_{l,g})^2 + \epsilon_{\text{probe}}}. \quad (16)$$

These probes describe rational curve shape on a shared interval; they are separate from the sampled live gains in equation 15. After the role/depth matrix step and immediately before reciprocal rescaling, the wrapper also measures the detached pair-state coordinate

$$\gamma_{l,g} = \log \text{rms}_{\epsilon_p}(A_{l,g}) - \log \text{rms}_{\epsilon_p}(B_{l,g}). \quad (17)$$

Only the probe gains and $\gamma_{l,g}$ are emitted by this summary family. The third control then computes $\tau_{l,g}$ and $\chi_{l,g}^{\text{rs}}$ from those summaries and from the pressure and activity memories.




Figure 2: Detached-summary route for the evaluated MatrixPolicy system. The top band produces stop-gradient summaries from forward hooks, pre-policy gradients, fixed-grid probes, and post-step pair-state measurements; the read-only bus routes them to three MatrixPolicy consumers. The lower policy band separates the gradient/update lane from the current matrix-pair state lane, while rational coefficients and non-RLB tensors remain AdamW-only.

3.2.1 GROUP-GRADIENT GATING

The first control redistributes gradient scale among groups inside each matrix role. It reads the sampled live gains, the pressure contrast $\pi_{l,g}$, and the activity contrast $\alpha_{l,g}$, then produces scaled gradients for A_l and B_l . A smooth gate $\phi_t = \text{smoothstep}(0.02, 0.30, \xi_t)$ turns on this control in the reported runs. The evaluated manifest sets gain, pressure, and activity strengths to 0.20,

0.10, and 0.20, and clips group multipliers to $[0.75, 1.35]$. For positive group values v_j , define $\text{geomean}_{j=1}^G(v_j) = \exp\left(G^{-1} \sum_{j=1}^G \log v_j\right)$. The group-gradient rule is

$$\begin{aligned}
 \tilde{\pi}_{l,g} &= \text{clip}(\pi_{l,g}, -1.50, 1.50), & c_{l,g}^{\text{activity}} &= \exp\left[-0.20\phi_t \text{ReLU}\left(\frac{\alpha_{l,g} - 0.05}{0.45}\right)\right], \\
 c_{l,g,\text{in}}^{\text{gain}} &= \left(\frac{\text{geomean}_{j=1}^G \hat{d}_{l,j}^{\text{live}}}{\hat{d}_{l,g}^{\text{live}}}\right)^{0.20\phi_t}, & c_{l,g,\text{out}}^{\text{gain}} &= \left(\frac{\text{geomean}_{j=1}^G \hat{o}_{l,j}^{\text{live}}}{\hat{o}_{l,g}^{\text{live}}}\right)^{0.20\phi_t}, \\
 c_{l,g,\text{in}}^{\text{pressure}} &= \exp(-0.10\phi_t \tilde{\pi}_{l,g}), & c_{l,g,\text{out}}^{\text{pressure}} &= \exp(+0.10\phi_t \tilde{\pi}_{l,g}), \\
 c_{l,g,\sigma}^{\text{raw}} &= c_{l,g,\sigma}^{\text{gain}} c_{l,g,\sigma}^{\text{pressure}} c_{l,g}^{\text{activity}}, & c_{l,g,\sigma} &= \text{clip}\left(\frac{c_{l,g,\sigma}^{\text{raw}}}{\text{geomean}_{j=1}^G c_{l,j,\sigma}^{\text{raw}}}, 0.75, 1.35\right).
 \end{aligned} \tag{18}$$

Live derivative gain controls the input role, live response gain controls the output role, and pressure acts with opposite signs on the two roles. When $\pi_{l,g} > 0$, the input-side relative-gradient EMA is larger than the output-side EMA, so the gate reduces the input-role gradient scale and increases the output-role gradient scale for that group. The activity term damps groups whose rational coefficient gradients are large relative to the adjacent matrix memories. The geometric recentering preserves the role-wise mean log multiplier before clipping; Appendix C.2 gives the log-space identity.

The scaled matrix gradients are

$$\tilde{\nabla} A_{l,g} = c_{l,g,\text{in}} \nabla A_{l,g}, \quad \tilde{\nabla} B_{l,g} = c_{l,g,\text{out}} \nabla B_{l,g}. \tag{19}$$

Let $\tilde{\nabla} M_{l,\text{in}}$ be the row-block concatenation of $\{\tilde{\nabla} A_{l,g}\}_{g=1}^G$, and let $\tilde{\nabla} M_{l,\text{out}}$ be the column-block concatenation of $\{\tilde{\nabla} B_{l,g}\}_{g=1}^G$.

3.2.2 ROLE/DEPTH MATRIX STEP

The second control applies the matrix optimizer step to the scaled gradients. It assigns different depth trends to the input and output roles, runs AdamW on both matrices, and activates a transient Muon substep early in training. In the reported configuration, the AdamW multiplier depends on role and depth; the pressure and activity memories enter this control through the Muon attenuation factor. Let $\delta_l = (l - 1)/(L - 1)$ for $L > 1$ and $\delta_l = 1/2$ otherwise. The branch schedules are

$$\begin{aligned}
 \varrho_{\text{in}}(l) &= \text{clip}(1 - 0.50(\delta_l - 0.5), 0.55, 1.40), \\
 \varrho_{\text{out}}(l) &= \text{clip}(1 + 1.00(\delta_l - 0.5), 0.55, 1.40), \\
 s_{l,\sigma}^{\text{Adam}} &= \text{clip}(3.0 \max\{0.10, 1 + 1.20(\varrho_{\sigma}(l) - 1)\}, 0.40, 4.0), \\
 \Psi_{l,t} &= \text{clip}\left(G^{-1} \sum_{g=1}^G |\pi_{l,g}|, 0, 1.50\right), \\
 \Omega_{l,t} &= \text{ReLU}\left(\frac{G^{-1} \sum_{g=1}^G \alpha_{l,g} - 0.05}{0.45}\right), \\
 \psi_{l,t} &= \text{clip}(\exp[-0.30\Psi_{l,t} - 0.65\Omega_{l,t}], 0.10, 1.15), \\
 \omega_t^\mu &= \text{smoothstep}(0.02, 0.12, \xi_t)[1 - \text{smoothstep}(0.20, 0.36, \xi_t)], \\
 \mu_{l,\sigma,t} &= \text{clip}(0.75 \omega_t^\mu \varrho_{\sigma}(l) \psi_{l,t}, 0, 0.75).
 \end{aligned} \tag{20}$$

Here $s_{l,\sigma}^{\text{Adam}}$ is the role/depth AdamW multiplier. The aggregate pressure $\Psi_{l,t}$ and activity $\Omega_{l,t}$ form the Muon attenuation $\psi_{l,t}$, so the active Muon fraction $\mu_{l,\sigma,t}$ is largest during the early Muon window and shrinks when the detached summaries indicate large role imbalance or large rational-coefficient activity. For $M_{l,\text{in}} = A_l$ and $M_{l,\text{out}} = B_l$, let η_t be the base learning rate and let $s_\mu = 1.0$ be the evaluated Muon learning-rate scale. MatrixPolicy applies an AdamW matrix step followed by any active Muon substep. The gate $\mu_{l,\sigma,t}$ scales the two optimizer learning rates; AdamW and Muon

keep separate states and are applied sequentially:

$$\begin{aligned}\widetilde{M}_{l,\sigma}^{t+1} &= \text{AdamW}_{\eta_t s_{l,\sigma}^{\text{Adam}}(1-\mu_{l,\sigma,t})}(M_{l,\sigma}^t; \widetilde{\nabla} M_{l,\sigma}^t), \\ M_{l,\sigma}^{t+1} &= \text{Muon}_{\eta_t s_{\mu} \mu_{l,\sigma,t}}(\widetilde{M}_{l,\sigma}^{t+1}; \widetilde{\nabla} M_{l,\sigma}^t),\end{aligned}\tag{21}$$

The AdamW step uses the training-protocol betas, weight decay, and learning-rate schedule in Appendix A.1. The Muon substep is the `torch.optim.Muon` update on two-dimensional RLB matrices with momentum 0.95, five Newton–Schulz steps, and match-RMS learning-rate adjustment (Liu et al., 2025); it becomes inactive when all $\mu_{l,\sigma,t}$ decay to zero. Appendix C.3 states the matrix-spectrum geometry behind this transient branch.

3.2.3 RECIPROCAL PAIR RESCALING

The third control chooses a bounded positive representative of the same RLB matrix pair after the matrix update. It is implemented by the surrounding transport wrapper and is evaluated every five optimizer steps after the AdamW/Muon branch; fixed-grid probe gains are refreshed every ten optimizer steps and reused between refreshes. It consumes the probe gains, the post-step log ratio $\gamma_{l,g}$, and the pressure/activity memories. The target combines the probe-gain balance coordinate derived in Appendix C.4 with a bounded pressure offset:

$$\begin{aligned}\bar{\pi}_{l,g} &= \text{clip}(\pi_{l,g}, -1.25, 1.25), \\ \tau_{l,g} &= \frac{1}{2} \left(\log \hat{d}_{l,g}^{\text{probe}} - \log \hat{o}_{l,g}^{\text{probe}} \right) + 0.25 \bar{\pi}_{l,g}, \\ \bar{\alpha}_{l,g} &= \text{clip}(\alpha_{l,g}, -1.25, 1.25), \\ \chi_{l,g}^{\text{rs}} &= \exp(0.10 \bar{\alpha}_{l,g}), \quad (\text{rescale-speed factor}) \\ s_{l,t} &= 0.50 \text{ smoothstep}(0.03, 0.35, \xi_t) \text{ clip}(1 + 0.15(\delta_l - 0.5), 0.75, 1.25).\end{aligned}\tag{22}$$

The current log ratio $\gamma_{l,g}$ is measured after the matrix step. The clipped log step and reciprocal update are

$$\begin{aligned}\ell_{l,g,t} &= \text{clip}\left(\frac{1}{2}[\tau_{l,g} - \gamma_{l,g}]s_{l,t}\chi_{l,g}^{\text{rs}}, -0.030, 0.030\right), \\ A_{l,g} &\leftarrow e^{\ell_{l,g,t}} A_{l,g}, \quad B_{l,g} \leftarrow e^{-\ell_{l,g,t}} B_{l,g}.\end{aligned}\tag{23}$$

The first term of $\tau_{l,g}$ equalizes derivative and response probe gains in the positive log coordinate, the pressure term shifts the target within a bounded range, and χ^{rs} changes the clipped rescale speed. Larger $\alpha_{l,g}$ accelerates this representative-balancing move, while the earlier gradient gate and Muon attenuation use the same activity contrast to reduce gradient-driven matrix motion. In the nonzero, unfloored idealization of equation 9, the reciprocal pair move preserves the represented function; near the RMS floor it is a clipped approximate balance step.

4 EXPERIMENTS

4.1 FIXED-SCALE TARGET-ARRIVAL STUDY

The fixed-scale study uses the same 12-layer, width-768 causal Transformer on five corpora: DCLM/DataComp-LM (Li et al., 2024), FineWeb-Edu and FineWeb (Penedo et al., 2024), Dolma-sample (Soldaini et al., 2024), and C4 (Raffel et al., 2020). The 100M-token budget uses 3050 optimizer steps and the 300M-token budget uses 9150 optimizer steps, with three seeds per method/dataset pair. The main comparison keeps the architecture and data fixed while changing the activation and optimizer used in the nonlinear sublayer: MatrixPolicy is compared with SiLU controls and with RLB-only controls, where RLB-only places the same activation under the named standard optimizer. Shared optimizer, evaluation-cadence, and throughput-measurement details are reported in Appendix A.1; broader optimizer comparisons beyond AdamW and Muon are summarized in Table 2.

Table 1: Target-arrival efficiency for the fixed 12-layer, width-768 Transformer. RLB denotes the same rational latent-basis activation without MatrixPolicy. At the hardest validation-loss target reached by every listed method in all three seeds, MatrixPolicy is always fastest; each token-saving column reports MatrixPolicy’s savings relative to the method named below it.

Budget	Dataset	Target loss	MatrixPolicy token savings (%)				Time to target (min)				
			vs. SiLU AdamW	vs. SiLU Muon	vs. RLB AdamW	vs. RLB Muon	Matrix Policy	SiLU AdamW	SiLU Muon	RLB AdamW	RLB Muon
100M	DCLM	4.55	20.7	32.4	20.7	34.3	12.8	20.8	26.0	14.3	19.0
	FineWeb-Edu	4.40	19.5	29.5	20.2	29.5	12.9	20.3	24.4	14.2	18.0
	FineWeb	4.60	22.8	32.1	21.5	32.1	12.6	16.4	20.2	16.2	20.1
	Dolma	4.60	22.0	34.0	22.7	34.9	12.8	17.6	22.0	17.3	22.2
	C4	4.60	20.7	29.2	19.3	28.7	11.6	14.3	17.3	14.9	18.8
300M	DCLM	4.10	23.3	8.3	24.0	7.0	39.6	53.5	45.9	54.7	45.6
	FineWeb-Edu	3.85	22.8	6.7	22.6	4.8	40.0	59.9	47.9	52.3	50.0
	FineWeb	4.10	23.1	7.1	22.4	5.3	40.8	61.2	46.4	55.6	46.6
	Dolma	3.95	22.3	9.8	22.3	6.5	44.4	49.3	52.3	50.5	47.4
	C4	4.00	26.4	8.5	25.3	6.5	41.8	64.7	55.5	59.4	51.9

Table 1 is the primary empirical result. It asks: for the hardest common validation-loss target reached by every listed method in all three seeds, what fraction of target-arrival tokens does MatrixPolicy save, and how many measured target-arrival minutes does each method require? Across all five datasets under the 100M-token budget, MatrixPolicy saves 19.5–22.8% of the SiLU with AdamW target-arrival tokens and 29.2–34.0% of the SiLU with Muon target-arrival tokens. Under the 300M-token budget, after longer training compresses endpoint gaps, MatrixPolicy still saves 22.3–26.4% versus SiLU with AdamW and 6.7–9.8% versus SiLU with Muon, with lower target-arrival time in every listed comparison.




Figure 3: MatrixPolicy shifts target arrival left across the completed fixed-scale study. Rows show the 100M- and 300M-token budgets, and columns show datasets. Each point is the mean token count needed to reach one validation-loss target in all three seeds; connected points trace one method from easier to harder targets. Black circles are MatrixPolicy. For every optimizer color, solid squares are RLB with that optimizer and dashed triangles are SiLU with the same optimizer. Lower targets are harder, and the hard-target timing and token-saving percentages are quantified in Table 1.




Figure 4: Representative 100M-token-budget training dynamics on DCLM and FineWeb-Edu. Rows show datasets and columns show validation loss, validation perplexity, and training loss. Lines are seed means over three runs and bands show one sample standard deviation. This figure uses the two standard SiLU baselines so that the early target-arrival effect in Table 1 remains visually legible; inset windows zoom the final-quarter saturation region.

The 100M-token-budget curves show why target arrival is the right main readout. MatrixPolicy separates early in validation loss and validation perplexity while keeping training loss competitive. A long fixed budget eventually compresses many activation/optimizer endpoints, especially under the 300M-token budget, so the paper emphasizes the number of tokens and minutes needed to reach the same useful validation-loss levels.

4.2 TODO: LARGER-SCALE TRAINING RESULTS

TODO: Insert the larger-scale training result here.

4.3 TODO: COMPONENT ABLATIONS

TODO: Insert the component ablation study here.

5 CONCLUSION

RLB turns the nonlinear Transformer sublayer into an optimizer-visible interface by exposing grouped normalized coordinates, trainable groupwise rational responses, and an idealized positive matrix-pair relation used as a bounded balancing coordinate. MatrixPolicy uses that interface to update the RLB input and output matrices with role-aware statistics, while rational coefficients and non-RLB-matrix parameters use AdamW. The completed 100M-token and 300M-token runs show that MatrixPolicy reaches common validation losses with fewer tokens and lower measured target-arrival time than the listed SiLU and RLB controls. [todo: Complete this conclusion after the larger-scale and ablation results are added.]

REFERENCES

- Nicolas Boullé, Yuji Nakatsukasa, and Alex Townsend. Rational Neural Networks. In *Advances in Neural Information Processing Systems*, 2020. URL <https://proceedings.neurips.cc/paper/2020/hash/a3f390d88e4c41f2747bfa2f1b5f87db-Abstract.html>.
- Xiangning Chen, Chen Liang, Da Huang, Esteban Real, Kaiyuan Wang, Hieu Pham, Xuanyi Dong, Thang Lu, Cho-Jui Hsieh, Yifeng Lu, and Quoc V. Le. Symbolic Discov-

- ery of Optimization Algorithms. In *Advances in Neural Information Processing Systems*, 2023. URL https://papers.neurips.cc/paper_files/paper/2023/hash/9a39b4925e35cf447ccba8757137d84f-Abstract-Conference.html.
- Aakanksha Chowdhery, Sharan Narang, Jacob Devlin, Maarten Bosma, Gaurav Mishra, Adam Roberts, Paul Barham, Hyung Won Chung, et al. PaLM: Scaling Language Modeling with Pathways. *arXiv preprint arXiv:2204.02311*, 2022. URL <https://arxiv.org/abs/2204.02311>.
- Yann N. Dauphin, Angela Fan, Michael Auli, and David Grangier. Language Modeling with Gated Convolutional Networks. In *Proceedings of the 34th International Conference on Machine Learning*, pp. 933–941, 2017. URL <https://proceedings.mlr.press/v70/dauphin17a.html>.
- Aaron Defazio, Xingyu Yang, Ahmed Khaled, Konstantin Mishchenko, Harsh Mehta, and Ashok Cutkosky. The Road Less Scheduled. In *Advances in Neural Information Processing Systems*, 2024. URL <https://neurips.cc/virtual/2024/poster/96925>.
- Haishuo Fang, Ji-Ung Lee, Nafise Sadat Moosavi, and Iryna Gurevych. Transformers with Learnable Activation Functions. In *Findings of the Association for Computational Linguistics: EACL 2023*, 2023. URL <https://aclanthology.org/2023.findings-eacl.181/>.
- Vineet Gupta, Tomer Koren, and Yoram Singer. Shampoo: Preconditioned Stochastic Tensor Optimization. In *Proceedings of the 35th International Conference on Machine Learning*, pp. 1842–1850, 2018. URL <https://proceedings.mlr.press/v80/gupta18a.html>.
- Dan Hendrycks and Kevin Gimpel. Gaussian Error Linear Units (GELUs). *arXiv preprint arXiv:1606.08415*, 2016. URL <https://arxiv.org/abs/1606.08415>.
- Diederik P. Kingma and Jimmy Ba. Adam: A Method for Stochastic Optimization. In *International Conference on Learning Representations*, 2015. URL <https://mlanthology.org/iclr/2015/kingma2015iclr-adam/>.
- Jeffrey Li, Alex Fang, Georgios Smyrnis, Maor Ivgi, Matt Jordan, Samir Gadre, Hritik Bansal, Etash Guha, et al. DataComp-LM: In Search of the Next Generation of Training Sets for Language Models. In *Advances in Neural Information Processing Systems, Datasets and Benchmarks Track*, 2024. doi: 10.52202/079017-0455. URL https://proceedings.neurips.cc/paper_files/paper/2024/hash/19e4ea30dded58259665db375885e412-Abstract-Datasets_and_Benchmarks_Track.html.
- Kaizhao Liang, Lizhang Chen, Bo Liu, and Qiang Liu. Cautious Optimizers: Improving Training with One Line of Code. In *International Conference on Learning Representations*, 2026. URL <https://openreview.net/forum?id=zBPZeRjfgu>.
- Hong Liu, Zhiyuan Li, David Hall, Percy Liang, and Tengyu Ma. Sophia: A Scalable Stochastic Second-order Optimizer for Language Model Pre-training. In *International Conference on Learning Representations*, 2024. URL https://proceedings.iclr.cc/paper_files/paper/2024/hash/06960915ba8674c7a898ec0b472b80ff-Abstract-Conference.html.
- Jingyuan Liu, Jianlin Su, Xingcheng Yao, Zhejun Jiang, Guokun Lai, Yulun Du, Yidao Qin, Weixin Xu, Enzhe Lu, Junjie Yan, Yanru Chen, Huabin Zheng, Yibo Liu, Shaowei Liu, Bohong Yin, Weiran He, Han Zhu, Yuzhi Wang, Jianzhou Wang, Mengnan Dong, Zheng Zhang, Yongsheng Kang, Hao Zhang, Xinran Xu, Yutao Zhang, Yuxin Wu, Xinyu Zhou, and Zhilin Yang. Muon is Scalable for LLM Training. *arXiv preprint arXiv:2502.16982*, 2025. URL <https://arxiv.org/abs/2502.16982>.
- Ilya Loshchilov and Frank Hutter. Decoupled Weight Decay Regularization. In *International Conference on Learning Representations*, 2019. URL <https://openreview.net/forum?id=Bkg6RiCqY7>.

- Yang Luo, Xiaozhe Ren, Zangwei Zheng, Zhuo Jiang, Xin Jiang, and Yang You. CAME: Confidence-guided Adaptive Memory Efficient Optimization. In *Proceedings of the 61st Annual Meeting of the Association for Computational Linguistics*, pp. 4442–4453, 2023. doi: 10.18653/v1/2023.acl-long.243. URL <https://aclanthology.org/2023.acl-long.243/>.
- Alejandro Molina, Patrick Schramowski, and Kristian Kersting. Padé Activation Units: End-to-End Learning of Flexible Activation Functions in Deep Networks. In *International Conference on Learning Representations*, 2020. URL <https://openreview.net/forum?id=BJlBSkHtDS>.
- Guilherme Penedo, Hynek Kydlicek, Loubna Ben Allal, Anton Lozhkov, Margaret Mitchell, Colin Raffel, Leandro von Werra, and Thomas Wolf. The FineWeb Datasets: Decanting the Web for the Finest Text Data at Scale. In *Advances in Neural Information Processing Systems, Datasets and Benchmarks Track*, 2024. doi: 10.52202/079017-0970. URL https://papers.nips.cc/paper/2024/hash/370df50ccfd8bde18f8f9c2d9151bda-Abstract-Datasets_and_Benchmarks_Track.html.
- Colin Raffel, Noam Shazeer, Adam Roberts, Katherine Lee, Sharan Narang, Michael Matena, Yanqi Zhou, Wei Li, and Peter J. Liu. Exploring the Limits of Transfer Learning with a Unified Text-to-Text Transformer. *Journal of Machine Learning Research*, 21(140):1–67, 2020. URL <http://jmlr.org/papers/v21/20-074.html>.
- Prajit Ramachandran, Barret Zoph, and Quoc V. Le. Searching for Activation Functions. *arXiv preprint arXiv:1710.05941*, 2017. URL <https://arxiv.org/abs/1710.05941>.
- Joel Schlotthauer, Christian Kroos, Chris Hinze, Viktor Hangya, Luzian Hahn, and Fabian Küch. Pre-Training LLMs on a Budget: A Comparison of Three Optimizers. *arXiv preprint arXiv:2507.08472*, 2025. URL <https://arxiv.org/abs/2507.08472>.
- Andrei Semenov, Matteo Pagliardini, and Martin Jaggi. Benchmarking Optimizers for Large Language Model Pretraining. *arXiv preprint arXiv:2509.01440*, 2025. URL <https://arxiv.org/abs/2509.01440>.
- Ishaan Shah, Anthony M. Polloreno, Karl Stratos, Philip Monk, Adarsh Chaluvareja, Andrew Hojel, Andrew Ma, Anil Thomas, Ashish Tanwer, Darsh J. Shah, Khoi Nguyen, Kurt Smith, Michael Callahan, Michael Pust, Mohit Parmar, Peter Rushton, Platon Mazarakis, Ritvik Kapila, Saurabh Srivastava, Somanshu Singla, Tim Romanski, Yash Vanjani, and Ashish Vaswani. Practical Efficiency of Muon for Pretraining. *arXiv preprint arXiv:2505.02222*, 2025. URL <https://arxiv.org/abs/2505.02222>.
- Noam Shazeer. GLU Variants Improve Transformer. *arXiv preprint arXiv:2002.05202*, 2020. URL <https://arxiv.org/abs/2002.05202>.
- Noam Shazeer and Mitchell Stern. Adafactor: Adaptive Learning Rates with Sublinear Memory Cost. In *Proceedings of the 35th International Conference on Machine Learning*, pp. 4596–4604, 2018. URL <https://proceedings.mlr.press/v80/shazeer18a.html>.
- Luca Soldaini, Rodney Kinney, Akshita Bhagia, Dustin Schwenk, David Atkinson, Russell Authur, Ben Bogin, Khyathi Chandu, et al. Dolma: An Open Corpus of Three Trillion Tokens for Language Model Pretraining Research. *arXiv preprint arXiv:2402.00159*, 2024. URL <https://arxiv.org/abs/2402.00159>.
- Jianlin Su, Yu Lu, Shengfeng Pan, Ahmed Murtadha, Bo Wen, and Yunfeng Liu. RoFormer: Enhanced Transformer with Rotary Position Embedding. *arXiv preprint arXiv:2104.09864*, 2021. URL <https://arxiv.org/abs/2104.09864>.
- Maosen Tang and Alex Townsend. Rational Neural Networks have Expressivity Advantages. *arXiv preprint arXiv:2602.12390*, 2026. doi: 10.48550/arXiv.2602.12390. URL <https://arxiv.org/abs/2602.12390>.

- Hugo Touvron, Thibaut Lavril, Gautier Izacard, Xavier Martinet, Marie-Anne Lachaux, Timothée Lacroix, Baptiste Rozière, Naman Goyal, et al. LLaMA: Open and Efficient Foundation Language Models. *arXiv preprint arXiv:2302.13971*, 2023. URL <https://arxiv.org/abs/2302.13971>.
- Ashish Vaswani, Noam Shazeer, Niki Parmar, Jakob Uszkoreit, Llion Jones, Aidan N. Gomez, Lukasz Kaiser, and Illia Polosukhin. Attention Is All You Need. In *Advances in Neural Information Processing Systems*, 2017. URL https://papers.neurips.cc/paper_files/paper/2017/hash/3f5ee243547dee91fbd053c1c4a845aa-Abstract.html.
- Nikhil Vyas, Depen Morwani, Rosie Zhao, Itai Shapira, David Brandfonbrener, Lucas Janson, and Sham Kakade. SOAP: Improving and Stabilizing Shampoo using Adam for Language Modeling. In *International Conference on Learning Representations*, 2025. URL https://proceedings.iclr.cc/paper_files/paper/2025/hash/e988664070e9591f93fddf605f7dc623-Abstract-Conference.html.
- Kaiyue Wen, David Leo Wright Hall, Tengyu Ma, and Percy Liang. Fantastic Pretraining Optimizers and Where to Find Them. In *International Conference on Learning Representations*, 2026. URL <https://openreview.net/forum?id=2J51qUZ0iG>.
- Biao Zhang and Rico Sennrich. Root Mean Square Layer Normalization. *arXiv preprint arXiv:1910.07467*, 2019. URL <https://arxiv.org/abs/1910.07467>.
- Yushun Zhang, Congliang Chen, Ziniu Li, Tian Ding, Chenwei Wu, Diederik P. Kingma, Yinyu Ye, Zhi-Quan Luo, and Ruoyu Sun. Adam-mini: Use Fewer Learning Rates To Gain More. In *International Conference on Learning Representations*, 2025. URL https://proceedings.iclr.cc/paper_files/paper/2025/hash/45ae878717399e6f62d57c65f052cd46-Abstract-Conference.html.
- Jiawei Zhao, Zhenyu Zhang, Beidi Chen, Zhangyang Wang, Anima Anandkumar, and Yuandong Tian. GaLore: Memory-Efficient LLM Training by Gradient Low-Rank Projection. In *International Conference on Machine Learning*, 2024. URL <https://openreview.net/forum?id=hYHsrKDIX7>.
- Rosie Zhao, Depen Morwani, David Brandfonbrener, Nikhil Vyas, and Sham Kakade. Deconstructing What Makes a Good Optimizer for Language Models. In *International Conference on Learning Representations*, 2025. URL <https://arxiv.org/abs/2407.07972>.

A EXPERIMENT APPENDIX

This appendix gives the fixed-scale protocol behind Section 4.1, the full training curves for the five corpora, and the broader optimizer endpoint comparison used to contextualize the main SiLU and RLB controls.

A.1 FIXED-SCALE TARGET-ARRIVAL STUDY

A.1.1 MODEL AND DATA PIPELINE

The fixed-scale experiments use a custom causal decoder Transformer. The block uses a LLaMA-style pre-normalized layout with RMSNorm before attention and before the nonlinear sublayer, bias-free query/key/value and output attention projections, rotary position embeddings, causal scaled-dot-product attention, residual attention and nonlinear-sublayer updates, a final RMSNorm, and a tied token-embedding/output head (Zhang & Sennrich, 2019; Su et al., 2021; Touvron et al., 2023). The fixed model has 12 layers, width 768, 12 heads of dimension 64, and context length 256.

The SiLU controls use a SwiGLU-like gated nonlinear sublayer with bias-free matrices $W_g, W_v \in \mathbb{R}^{768 \times 2048}$ and $W_o \in \mathbb{R}^{2048 \times 768}$, computing $(\text{SiLU}(xW_g) \odot xW_v)W_o$. MatrixPolicy and RLB-only configurations use a matched single-branch matrix pair. An input matrix maps width 768 to 3072 latent coordinates, the coordinates are partitioned into groups and transformed by the real-pole-free P5/Q4 rational response, and an output matrix maps the 3072 coordinates back to width

768. With the default group size 256, this gives $G = 12$ groups of width $m = 256$; the RMS floor is $\epsilon_{\text{rms}} = 10^{-6}$. This is the RLB configuration used by the main MatrixPolicy and RLB-control results.

Tokenization uses the GPT-2 tokenizer through the Hugging Face tokenizer interface with its base vocabulary; an EOS or pad token is appended between documents. Training samples random contiguous blocks of length 257 and forms next-token prediction pairs of length 256. The DCLM, FineWeb-Edu, FineWeb, Dolma-sample, and C4 runs use the text field specified in the manifests, with validation slices drawn from held-out token offsets of the same corpus stream.

A.1.2 TRAINING PROTOCOL

All runs use batch size 16, gradient accumulation 2, and therefore 32768 global tokens per optimizer step. The 100M-token budget consists of 3050 optimizer steps; the 300M-token budget consists of 9150 optimizer steps. Each method/dataset pair is run with seeds 1337, 2027, and 3407. Validation is performed every 50 optimizer steps, so token-to-target arrivals are quantized in increments of 1.64M tokens.

The MatrixPolicy rows in the main table use the RLB activation, whose active nonlinearity is the groupwise P5/Q4 response in equations 3–5. The evaluated configuration uses AdamW as the backbone optimizer and disables the separate rational-coefficient optimizer. It enables MatrixPolicy group redistribution with gain/pressure/activity strengths 0.20/0.10/0.20, schedule window 0.02–0.30, and group multiplier clip $[0.75, 1.35]$. The RLB matrix pair then uses the role/depth-scaled AdamW branch and the early Muon substep from Section 3.2. The bounded RLB pair-rescale wrapper runs every five optimizer steps after the matrix branch. Rational coefficients, attention weights, embeddings, normalization parameters, and ordinary Transformer weights use AdamW in this configuration.

Target-arrival minutes in Table 1 are computed seed-wise from the first validation event that reaches the shared target and the same run’s measured seconds per optimizer step. The per-step time is the training-script wall time recorded in the run summary divided by completed optimizer steps for that run; target arrivals are therefore measured with the same evaluation cadence and timing source across methods.

A.1.3 SUPPORTING FIXED-SCALE CURVES




Figure 5: 100M-token-budget dynamics across all five datasets. Rows are datasets and columns are validation loss, validation perplexity, and training loss. Lines are seed means over three runs and bands show one sample standard deviation. Insets zoom the final-quarter saturation region in each panel.




Figure 6: 300M-token-budget dynamics across all five datasets. Rows are datasets and columns are validation loss, validation perplexity, and training loss. Lines are seed means over three runs and bands show one sample standard deviation. Insets zoom the middle training window where early saturation differences are most visible.

A.1.4 BROAD OPTIMIZER ENDPOINT TABLE

The fixed-scale package also includes the broader optimizer endpoints referenced from Section 4.1. These rows use the same model, datasets, seeds, token budgets, and validation schedule as the main comparison, while replacing the nonlinear-sublayer optimizer according to the table entry.

Table 2: Final validation loss for the broader optimizer sweep on the fixed 12-layer, width-768 language model. Values are mean $\pm$ sample standard deviation over three seeds; lower is better. RLB rows use the same rational latent-basis activation without MatrixPolicy.

Budget	Method	DCLM	FineWeb-Edu	FineWeb	Dolma	C4
100M tokens	MatrixPolicy	4.254 $\pm$ 0.006	4.087 $\pm$ 0.010	4.316 $\pm$ 0.013	4.325 $\pm$ 0.005	4.284 $\pm$ 0.020
	SiLU + AdamW	4.406 $\pm$ 0.010	4.237 $\pm$ 0.009	4.476 $\pm$ 0.010	4.486 $\pm$ 0.001	4.447 $\pm$ 0.016
	SiLU + Lion	4.318 $\pm$ 0.007	4.149 $\pm$ 0.009	4.383 $\pm$ 0.008	4.388 $\pm$ 0.005	4.354 $\pm$ 0.016
	SiLU + Muon	4.457 $\pm$ 0.013	4.279 $\pm$ 0.024	4.516 $\pm$ 0.026	4.544 $\pm$ 0.011	4.482 $\pm$ 0.023
	SiLU + SOAP	4.416 $\pm$ 0.004	4.263 $\pm$ 0.022	4.484 $\pm$ 0.011	4.499 $\pm$ 0.004	4.458 $\pm$ 0.015
	SiLU + ScheduleFree	4.902 $\pm$ 0.011	4.826 $\pm$ 0.007	5.014 $\pm$ 0.019	5.049 $\pm$ 0.015	5.006 $\pm$ 0.016
	SiLU + CAME	5.011 $\pm$ 0.014	4.921 $\pm$ 0.012	5.133 $\pm$ 0.013	5.165 $\pm$ 0.019	5.123 $\pm$ 0.018
	RLB + AdamW	4.405 $\pm$ 0.008	4.239 $\pm$ 0.008	4.472 $\pm$ 0.014	4.488 $\pm$ 0.003	4.441 $\pm$ 0.016
	RLB + Lion	4.295 $\pm$ 0.008	4.136 $\pm$ 0.008	4.363 $\pm$ 0.011	4.362 $\pm$ 0.007	4.327 $\pm$ 0.016
	RLB + Muon	4.467 $\pm$ 0.012	4.283 $\pm$ 0.028	4.516 $\pm$ 0.006	4.548 $\pm$ 0.015	4.476 $\pm$ 0.010
	RLB + SOAP	4.436 $\pm$ 0.044	4.275 $\pm$ 0.024	4.689 $\pm$ 0.336	4.518 $\pm$ 0.036	4.446 $\pm$ 0.019
	RLB + ScheduleFree	4.887 $\pm$ 0.005	4.796 $\pm$ 0.013	4.999 $\pm$ 0.019	5.036 $\pm$ 0.013	4.984 $\pm$ 0.016
	RLB + CAME	5.014 $\pm$ 0.009	4.915 $\pm$ 0.006	5.134 $\pm$ 0.016	5.151 $\pm$ 0.017	5.111 $\pm$ 0.016
300M tokens	MatrixPolicy	3.952 $\pm$ 0.028	3.702 $\pm$ 0.021	3.962 $\pm$ 0.008	3.806 $\pm$ 0.007	3.878 $\pm$ 0.014
	SiLU + AdamW	4.049 $\pm$ 0.027	3.803 $\pm$ 0.018	4.061 $\pm$ 0.010	3.904 $\pm$ 0.009	3.981 $\pm$ 0.013
	SiLU + Lion	3.993 $\pm$ 0.023	3.744 $\pm$ 0.021	4.001 $\pm$ 0.008	3.848 $\pm$ 0.009	3.921 $\pm$ 0.011
	SiLU + Muon	3.997 $\pm$ 0.030	3.745 $\pm$ 0.017	4.007 $\pm$ 0.013	3.858 $\pm$ 0.010	3.925 $\pm$ 0.013
	SiLU + SOAP	4.096 $\pm$ 0.030	3.863 $\pm$ 0.020	4.114 $\pm$ 0.010	3.957 $\pm$ 0.009	4.035 $\pm$ 0.011
	SiLU + ScheduleFree	4.366 $\pm$ 0.030	4.156 $\pm$ 0.024	4.398 $\pm$ 0.011	4.215 $\pm$ 0.005	4.316 $\pm$ 0.011
	SiLU + CAME	4.368 $\pm$ 0.023	4.150 $\pm$ 0.021	4.406 $\pm$ 0.019	4.249 $\pm$ 0.027	4.330 $\pm$ 0.015
	RLB + AdamW	4.050 $\pm$ 0.030	3.802 $\pm$ 0.021	4.060 $\pm$ 0.011	3.905 $\pm$ 0.008	3.979 $\pm$ 0.010
	RLB + Lion	3.989 $\pm$ 0.029	3.742 $\pm$ 0.021	3.996 $\pm$ 0.011	3.841 $\pm$ 0.008	3.913 $\pm$ 0.014
	RLB + Muon	3.991 $\pm$ 0.026	3.737 $\pm$ 0.019	3.999 $\pm$ 0.011	3.848 $\pm$ 0.005	3.919 $\pm$ 0.015
	RLB + SOAP	4.061 $\pm$ 0.033	3.824 $\pm$ 0.013	4.108 $\pm$ 0.032	3.927 $\pm$ 0.013	4.002 $\pm$ 0.019
	RLB + ScheduleFree	4.361 $\pm$ 0.034	4.142 $\pm$ 0.022	4.386 $\pm$ 0.008	4.212 $\pm$ 0.011	4.308 $\pm$ 0.007
	RLB + CAME	4.450 $\pm$ 0.043	4.225 $\pm$ 0.036	4.480 $\pm$ 0.006	4.267 $\pm$ 0.041	4.365 $\pm$ 0.058

A.1.5 LEARNING-RATE AND WEIGHT-DECAY SENSITIVITY

We also ran a 100M-token, single-seed sensitivity grid for MatrixPolicy and the two standard SiLU controls. The grid crosses learning rates $\{1, 2, 3, 5\} \times 10^{-4}$ with weight decays 0, 0.05, 0.10, and 0.20, using the same model, corpora, tokenizer, evaluation cadence, and timing source as Appendix A.1.2. For each dataset, Table 3 anchors the analysis at the hard 100M-token target used in Table 1 and includes targets 0.10 and 0.20 loss units below it. Each row uses the paired learning-rate/weight-decay cells in which all three methods reach that target; $n/16$ reports the resulting common-arrival coverage. Target-arrival minutes equal the first-hit step multiplied by the measured run-level mean seconds per step.

Across all 15 dataset–target rows, MatrixPolicy has positive median token savings relative to both SiLU controls. At the hard main-study targets, the median saving is 14.4–18.1% relative to SiLU with AdamW and 18.8–22.1% relative to SiLU with Muon. Across the two lower targets, the corresponding ranges are 8.8–21.4% and 9.8–23.2%, respectively. MatrixPolicy has lower median target-arrival time than SiLU with Muon in all 15 rows and matches or improves the SiLU-with-AdamW median time in 12 rows. The target-arrival advantage therefore persists across the moderate learning-rate and weight-decay perturbations that reach these lower loss regimes.

Table 3: Single-seed target-arrival sensitivity at the 100M-token budget. Each dataset uses the hard target from the main 100M-token study and two progressively lower targets. Entries are medians over the common learning-rate/weight-decay cells in which all three methods reach the target. The paired-cells column reports common-arrival coverage out of 16.

Dataset	Target loss	Paired cells	MatrixPolicy token savings (%)		Time to target (min)		
			vs. SiLU AdamW	vs. SiLU Muon	Matrix Policy	SiLU AdamW	SiLU Muon
DCLM	4.55	8/16	14.4	19.1	11.1	11.1	13.1
	4.45	7/16	8.8	11.4	12.0	11.4	12.7
	4.35	4/16	11.9	10.8	13.7	13.8	15.2
FineWeb-Edu	4.40	8/16	14.4	19.1	10.8	10.8	12.7
	4.30	8/16	19.6	23.2	12.7	13.4	15.9
	4.20	4/16	14.0	9.8	13.7	13.8	14.6
FineWeb	4.60	8/16	16.4	19.1	10.7	11.0	12.9
	4.50	8/16	21.4	22.0	12.8	13.8	16.3
	4.40	4/16	15.9	10.8	14.2	14.3	14.7
Dolma	4.60	8/16	18.1	22.1	11.3	11.8	14.1
	4.50	4/16	13.5	13.5	12.5	12.3	13.3
	4.40	4/16	14.3	13.3	15.3	15.0	16.2
C4	4.60	8/16	14.8	18.8	10.4	10.6	12.3
	4.50	8/16	19.3	21.7	12.2	13.3	15.2
	4.40	4/16	14.6	10.3	13.0	13.3	13.9

A.2 TODO: LARGER-SCALE TRAINING RESULTS

TODO: Insert larger-scale training details here.

A.3 TODO: COMPONENT ABLATIONS

TODO: Insert component ablation details here.

B RATIONAL LATENT BASIS CALCULUS

The RLB summaries and the reciprocal matrix-pair coordinate follow from the scalar response and the RMS-normalized group map. We suppress layer and group indices throughout this section.

B.1 POLE-FREE RATIONAL RESPONSE

Let $\phi(u)$ collect the nonconstant denominator features. The evaluated P5/Q4 response is

$$\begin{aligned}
 P(u) &:= \sum_{i=0}^5 a_i u^i, & \phi(u) &:= (|u|, u^2, |u|^3, u^4), \\
 Q(u) &:= 1 + \sum_{j=1}^4 |b_j| \phi_j(u), & R(u) &:= \frac{P(u)}{Q(u)}.
 \end{aligned} \tag{24}$$

Because each denominator feature is nonnegative, $Q(u) \geq 1$ for every $u \in \mathbb{R}$; hence the response is pole-free on $\mathbb{R}$. Away from the kink at $u = 0$,

$$\begin{aligned}
 R'(u) &= \frac{P'(u)Q(u) - P(u)Q'(u)}{Q(u)^2}, \\
 P'(u) &= \sum_{i=1}^5 i a_i u^{i-1}, & Q'(u) &= |b_1| \operatorname{sign}(u) + 2|b_2|u + 3|b_3|u|u| + 4|b_4|u^3.
 \end{aligned} \tag{25}$$

The selected coefficient derivatives are

$$\frac{\partial R(u)}{\partial a_i} = \frac{u^i}{Q(u)}, \quad \frac{\partial R(u)}{\partial b_j} = -\frac{P(u)}{Q(u)^2} \text{sign}(b_j) \phi_j(u). \quad (26)$$

The implementation selects $\text{sign}(0) = 0$ for both $u = 0$ and $b_j = 0$. At these kinks, the formulas specify the backward derivative chosen by autodiff. For vector inputs, R and R' act coordinatewise. The response, derivative, and coefficient-activity summaries therefore originate from $R(u)$, $R'(u)$, and the features in equation 26, respectively. For the evaluated response, $\Theta_{l,g}^R = \{a_{l,g,\cdot}, b_{l,g,\cdot}\}$ in equation 12.

B.2 RMS-NORMALIZED GROUP DIFFERENTIAL

Let m be the group width and let $\epsilon_{\text{rms}} \geq 0$. For $\epsilon_{\text{rms}} > 0$ and $z \in \mathbb{R}^m$, or for $\epsilon_{\text{rms}} = 0$ and $z \neq 0$, define

$$r(z) = \sqrt{m^{-1}\|z\|_2^2 + \epsilon_{\text{rms}}}, \quad u(z) = \frac{z}{r(z)}, \quad h(z) = r(z)R(u(z)). \quad (27)$$

On the domain $r(z) > 0$, direct differentiation gives

$$dr = \frac{z^\top dz}{mr} = \frac{u^\top dz}{m}, \quad du = \frac{1}{r} \left(I_m - \frac{uu^\top}{m} \right) dz. \quad (28)$$

The factor $I_m - uu^\top/m$ is the exact normalized-coordinate differential. It is the tangent projector to the RMS sphere when $\epsilon_{\text{rms}} = 0$. With a positive floor, its radial eigenvalue is $\epsilon_{\text{rms}}/r^2$. Away from scalar-response kinks, the group Jacobian is

$$J_h(z) := \frac{\partial h}{\partial z} = R(u) \frac{u^\top}{m} + \text{diag}(R'(u)) \left(I_m - \frac{uu^\top}{m} \right). \quad (29)$$

This decomposition identifies the scalar factors observed by the live summaries. Response RMS summarizes the radial response factor, and derivative RMS summarizes $\text{diag}(R'(u))$. The complete local Jacobian additionally depends on the normalized-coordinate factor and radial term; the matrix-role gradient also depends on the output matrix and upstream adjoint. The two RMS summaries therefore serve as role-specific scalar proxies.

B.3 POSITIVE PAIR COORDINATE AND FLOOR PERTURBATION

Set $\epsilon_{\text{rms}} = 0$ and write r_0, u_0, h_0 for the unfloored map. For $z \neq 0$ and $\lambda > 0$,

$$r_0(\lambda z) = \lambda r_0(z), \quad u_0(\lambda z) = u_0(z), \quad h_0(\lambda z) = \lambda h_0(z). \quad (30)$$

For a group contribution $y_{l,g} = B_{l,g}h(A_{l,g}x_l)$, with R fixed, $\epsilon_{\text{rms}} = 0$, and $A_{l,g}x_l \neq 0$, the transformation $(A_{l,g}, B_{l,g}) \mapsto (\lambda A_{l,g}, \lambda^{-1} B_{l,g})$ leaves $y_{l,g}$ unchanged. This is the positive pair orbit in equation 9.

The RMS floor perturbs the orbit through both the radius and normalized coordinate. Write

$$r_\epsilon(z) = \sqrt{m^{-1}\|z\|_2^2 + \epsilon_{\text{rms}}}, \quad u_\epsilon(z) = z/r_\epsilon(z), \quad h_\epsilon(z) = r_\epsilon(z)R(u_\epsilon(z)). \quad (31)$$

For $\rho^2 = m^{-1}\|z\|_2^2 > 0$ and $\delta = \epsilon_{\text{rms}}/\rho^2$, define

$$s_\epsilon(\lambda, z) := \frac{r_\epsilon(\lambda z)}{\lambda r_\epsilon(z)} = \sqrt{\frac{1 + \delta/\lambda^2}{1 + \delta}}, \quad \kappa(\lambda, z) := s_\epsilon(\lambda, z)^{-1}, \quad (32)$$

$$u_\epsilon(\lambda z) = \kappa(\lambda, z)u_\epsilon(z).$$

Fix $c > 0$, and let $\mathcal{I}_c$ contain every scalar coordinate on all segments joining $u_\epsilon(z)$ and $\kappa(\lambda, z)u_\epsilon(z)$ for $\lambda \in [e^{-c}, e^c]$ and all sufficiently small δ . Suppose that R is L_R -Lipschitz on $\mathcal{I}_c$. Reciprocal output compensation then gives, for every λ in this interval, the latent feature bound

$$\begin{aligned} \|\lambda^{-1}h_\epsilon(\lambda z) - h_\epsilon(z)\|_2 &\leq \left| \frac{r_\epsilon(\lambda z)}{\lambda} - r_\epsilon(z) \right| \|R(u_\epsilon(z))\|_2 \\ &\quad + \frac{r_\epsilon(\lambda z)}{\lambda} L_R |\kappa(\lambda, z) - 1| \|u_\epsilon(z)\|_2. \end{aligned} \quad (33)$$

After the pair map, the one-group output difference is $B_{l,g}[\lambda^{-1}h_\epsilon(\lambda z) - h_\epsilon(z)]$, where $B_{l,g}$ denotes the pre-rescale output block. Its norm is at most $\|B_{l,g}\|_{\text{op}}$ times the right-hand side of equation 33. For fixed c and fixed m , uniformly over $\lambda \in [e^{-c}, e^c]$ as $\delta \rightarrow 0$,

$$\sup_{\lambda \in [e^{-c}, e^c]} (|s_\epsilon(\lambda, z) - 1| + |\kappa(\lambda, z) - 1|) = O_c(\delta),$$

$$\sup_{\lambda \in [e^{-c}, e^c]} \|\lambda^{-1}h_\epsilon(\lambda z) - h_\epsilon(z)\|_2 = O_c\left(\sqrt{m} [\sup_{\mathcal{I}_c} |R| + L_R] \frac{\epsilon_{\text{rms}}}{\rho}\right). \quad (34)$$

The first line follows directly from $|s_\epsilon - 1| \leq \frac{1}{2}\delta|\lambda^{-2} - 1|$ and $|\kappa - 1| = |s_\epsilon - 1|/s_\epsilon$. For radii separated from zero and bounded reciprocal moves, the error decays as $\epsilon_{\text{rms}}/\rho$; this rate motivates the clipped log-scale update in MatrixPolicy.

C OPTIMIZER DESIGN AND MATRIXPOLICY GEOMETRY

Let $\text{sg}(\cdot)$ denote stop-gradient. MatrixPolicy first forms detached role summaries, then rescales the current matrix gradients, applies the AdamW/Muon matrix branch, and finally moves the updated pair along the positive coordinate from Appendix B.3. Its effective order is

$$\begin{aligned} \nabla \mathcal{L} &\longrightarrow \text{sg}\left(q^{\text{in}}, q^{\text{out}}, q^R, \pi, \alpha, \hat{d}^{\text{live}}, \hat{o}^{\text{live}}\right) \longrightarrow (\tilde{\nabla} A, \tilde{\nabla} B) \\ &\longrightarrow (A^{\text{step}}, B^{\text{step}}) \longrightarrow \text{sg}\left(\hat{d}^{\text{probe}}, \hat{o}^{\text{probe}}, \gamma\right) \\ &\longrightarrow (e^\ell A^{\text{step}}, e^{-\ell} B^{\text{step}}). \end{aligned} \quad (35)$$

C.1 LOCAL MATRIX-ROLE GEOMETRY

For one token of one training example, let $\bar{y}_l = \nabla_{y_l} \mathcal{L}$ and $J_{l,g} = \partial h_{l,g} / \partial z_{l,g}$ from equation 29. Here x_l is the RLB input, $z_{l,g} = A_{l,g} x_l$, $h_{l,g} = h(z_{l,g})$, and $y_l = \sum_g B_{l,g} h_{l,g}$. The adjacent matrix gradients factor as

$$\nabla_{B_{l,g}} \mathcal{L} = \bar{y}_l h_{l,g}^\top, \quad \nabla_{A_{l,g}} \mathcal{L} = (J_{l,g}^\top B_{l,g}^\top \bar{y}_l) x_l^\top. \quad (36)$$

Batch and sequence gradients are sums or averages of these outer-product terms. The output-role gradient contains the realized feature $h_{l,g}$, whereas the input-role gradient contains the local Jacobian $J_{l,g}$. This factorization supports response RMS and derivative RMS as scalar role proxies. The full gradient retains the omitted radius, normalized-coordinate factor, matrix orientation, data distribution, and upstream adjoint.

The pressure summaries quantify relative gradient scale. For a pure gradient step $W^+ = W - \eta \mathcal{G}$, with $W \neq 0$ and $\eta \|\mathcal{G}\|_F \leq c \|W\|_F$ for fixed $c < 1$, Taylor expansion gives

$$\log \text{rms}(W^+) - \log \text{rms}(W) = -\eta \frac{\langle \mathcal{G}, W \rangle_F}{\|W\|_F^2} + O_c\left(\eta^2 \frac{\|\mathcal{G}\|_F^2}{\|W\|_F^2}\right), \quad (37)$$

while Cauchy–Schwarz yields

$$\left| \frac{\langle \mathcal{G}, W \rangle_F}{\|W\|_F^2} \right| \leq \frac{\|\mathcal{G}\|_F}{\|W\|_F} = \frac{\text{rms}(\mathcal{G})}{\text{rms}(W)}. \quad (38)$$

Thus $\text{rms}(\mathcal{G})/\text{rms}(W)$ bounds the magnitude of the first-order log-RMS change for a pure gradient step. The implemented p^{in} and p^{out} are floored versions of this nonnegative bound, computed from post-global-clipping gradients before the group policy. They summarize unsigned pre-policy gradient scale; AdamW and Muon subsequently construct signed, state-dependent updates. Their EMA contrast $\pi_{l,g}$ identifies the role with larger recent relative gradient scale. The activity contrast $\alpha_{l,g}$ compares raw rational-coefficient gradient energy with the geometric mean of the matrix-pressure memories and supplies a calibrated policy signal for relative coefficient activity.

In the unfloored fixed-response map, the positive pair orbit also satisfies

$$\left. \frac{d}{ds} \mathcal{L}(e^s A_{l,g}, e^{-s} B_{l,g}) \right|_{s=0} = \langle \nabla A_{l,g}, A_{l,g} \rangle_F - \langle \nabla B_{l,g}, B_{l,g} \rangle_F = 0. \quad (39)$$

The final control follows this zero-loss tangent in the idealized map to alter relative matrix pressure.

C.2 CENTERED GROUP CONTROL

The first control is most transparent in log space. Define the layerwise mean log gains

$$\bar{d}_l = \frac{1}{G} \sum_{j=1}^G \log \hat{d}_{l,j}^{\text{live}}, \quad \bar{o}_l = \frac{1}{G} \sum_{j=1}^G \log \hat{o}_{l,j}^{\text{live}}. \quad (40)$$

The gain factors in equation 18 obey

$$\log c_{l,g,\text{in}}^{\text{gain}} = -0.20\phi_t(\log \hat{d}_{l,g}^{\text{live}} - \bar{d}_l), \quad \log c_{l,g,\text{out}}^{\text{gain}} = -0.20\phi_t(\log \hat{o}_{l,g}^{\text{live}} - \bar{o}_l). \quad (41)$$

Above-geometric-mean derivative or response gain therefore reduces the corresponding role multiplier. The pressure factors act with opposite signs,

$$\log c_{l,g,\text{in}}^{\text{pressure}} = -0.10\phi_t \tilde{\pi}_{l,g}, \quad \log c_{l,g,\text{out}}^{\text{pressure}} = +0.10\phi_t \tilde{\pi}_{l,g}, \quad (42)$$

so positive $\pi_{l,g}$ reduces the input-role gradient and increases the output-role gradient before recentering. The shared activity factor is

$$\log c_{l,g}^{\text{activity}} = -0.20\phi_t \text{ReLU}\left(\frac{\alpha_{l,g} - 0.05}{0.45}\right). \quad (43)$$

To distinguish the stages of the implementation, let

$$c_{l,g,\sigma}^{\text{pre}} := c_{l,g,\sigma}^{\text{gain}} c_{l,g,\sigma}^{\text{pressure}} c_{l,g}^{\text{activity}},$$

$$\hat{c}_{l,g,\sigma} := \frac{c_{l,g,\sigma}^{\text{pre}}}{\text{geomean}_{j=1}^G c_{l,j,\sigma}^{\text{pre}}}, \quad c_{l,g,\sigma} := \text{clip}(\hat{c}_{l,g,\sigma}, 0.75, 1.35), \quad (44)$$

$$\frac{1}{G} \sum_{g=1}^G \log \hat{c}_{l,g,\sigma} = 0.$$

The zero-mean identity holds exactly before the final clip; the clip may shift the layerwise mean log multiplier. The gate preserves each group-gradient direction and multiplies its magnitude by a positive scalar before AdamW and Muon construct their state-dependent update directions.

C.3 TRANSIENT SPECTRAL NORMALIZATION

The role/depth factors in equation 20 are bounded scalar preconditioners. They preserve each matrix direction while changing the step magnitude assigned to the input and output maps across depth. AdamW then maintains entrywise moments (Kingma & Ba, 2015; Loshchilov & Hutter, 2019),

$$M \leftarrow \beta_1 M + (1 - \beta_1)G, \quad V \leftarrow \beta_2 V + (1 - \beta_2)(G \odot G),$$

$$D_{ij}^{\text{Adam}} = \frac{M_{ij}}{\sqrt{V_{ij}} + \epsilon}, \quad (45)$$

where the display suppresses bias correction and decoupled weight decay. This normalization acts coordinatewise; the subsequent Muon branch controls the singular spectrum of the matrix direction.

Matrix spectra matter for the RLB pair because equation 36 gives single-example adjacent gradients as outer products; after averaging over batch and sequence positions, the training gradients are sums of such terms. Singular modes of B_l couple output-residual directions to latent features, and singular modes of A_l couple latent coordinates to input-residual directions. During early basis formation, a highly anisotropic singular spectrum concentrates movement in a narrow set of these coupled directions.

Let $H_{l,\sigma,t} = U\Sigma V^\top$ be the compact SVD of the tensor passed to Newton–Schulz after group-gradient scaling and Muon’s momentum/Nesterov construction. Its ideal polar direction is

$$\mathcal{P}(H_{l,\sigma,t}) = UV^\top, H = \sum_i \sigma_i u_i v_i^\top \mapsto \sum_i u_i v_i^\top \quad \text{on the nonzero singular subspace.} \quad (46)$$

The polar map preserves the row and column singular vectors, replaces every nonzero singular value by one, and retains the nonzero singular subspace of $H_{l,\sigma,t}$. The polar factor solves

$$\mathcal{P}(H) \in \arg \max_{\|D\|_2 \leq 1} \langle D, H \rangle_F, \quad \langle \mathcal{P}(H), H \rangle_F = \|H\|_*, \quad (47)$$

$$\mathcal{P}(aH) = \mathcal{P}(H) \quad (a > 0, H \neq 0).$$

The implementation uses $D_{l,\sigma,t}^\mu = \text{NS}_5(H_{l,\sigma,t})$, a finite Newton–Schulz approximation to the polar direction, together with match-RMS learning-rate adjustment (Liu et al., 2025).

The branch equation in equation 21 is sequential because AdamW and Muon keep separate states. Let $\mathcal{D}_{l,\sigma,t}^{\text{Adam}}$ and $\mathcal{D}_{l,\sigma,t}^\mu$ denote their complete first-order directions, including bias correction, decoupled weight decay, and, for Muon, the finite Newton–Schulz and match-RMS adjustments. The sequential step has expansion

$$\Delta M_{l,\sigma} = -\eta_t s_{l,\sigma}^{\text{Adam}} (1 - \mu_{l,\sigma,t}) \mathcal{D}_{l,\sigma,t}^{\text{Adam}} - \eta_t s_\mu \mu_{l,\sigma,t} \mathcal{D}_{l,\sigma,t}^\mu + O(\eta_t^2), \quad (48)$$

after the group control has set the role-specific gradient scales. Flattening the nonzero singular values gives weaker represented modes more relative movement within the existing singular subspace. MatrixPolicy confines this intervention to the early window in equation 20; AdamW and reciprocal pair balancing remain after the Muon fraction reaches zero.

The Muon gate in equation 20 is monotone in the two layer summaries. Away from clipping boundaries and the ReLU kink, the derivatives of $\log \psi_{l,t}$ with respect to $\Psi_{l,t}$ and $\Omega_{l,t}$ are -0.30 and -0.65 inside the ψ clipping interval. Thus the Muon fraction is largest when role pressure and rational-coefficient activity are both small. Large $|\pi_{l,g}|$ marks unequal input/output pressure, and large $\alpha_{l,g}$ marks rational coefficients moving strongly relative to the matrix pair. These are precisely the regimes where the spectrum-normalized substep is attenuated.

C.4 RECIPROCAL PAIR BALANCING

After the descent branch, MatrixPolicy applies the reciprocal reparameterization

$$A_{l,g}^+ = e^{\ell_{l,g,t}} A_{l,g}, \quad B_{l,g}^+ = e^{-\ell_{l,g,t}} B_{l,g}, \quad (49)$$

every five optimizer steps. The construction follows the positive coordinate in Appendix B.3. With the rational response held fixed, the unfloored map satisfies $h_{l,g}(\lambda A_{l,g} x) = \lambda h_{l,g}(A_{l,g} x)$ for $\lambda > 0$, so replacing $(A_{l,g}, B_{l,g})$ by $(\lambda A_{l,g}, \lambda^{-1} B_{l,g})$ leaves $B_{l,g} h_{l,g}$ unchanged. The floor-error bound in equation 33 quantifies the departure from this identity.

Equivalent unfloored representatives have different relative gradient pressures. Let

$$p_{l,g}^A = \frac{\text{rms}(\nabla A_{l,g})}{\text{rms}(A_{l,g})}, \quad p_{l,g}^B = \frac{\text{rms}(\nabla B_{l,g})}{\text{rms}(B_{l,g})}. \quad (50)$$

In the fixed-response, unfloored map, $J_h(e^\ell z) = J_h(z)$ almost everywhere. Evaluating the instantaneous gradients at the transformed representative gives

$$\begin{aligned} \nabla_{A_{l,g}^+} \mathcal{L} &= e^{-\ell_{l,g,t}} \nabla_{A_{l,g}} \mathcal{L}, & \nabla_{B_{l,g}^+} \mathcal{L} &= e^{\ell_{l,g,t}} \nabla_{B_{l,g}} \mathcal{L}, \\ p_{l,g}^{A,+} &= e^{-2\ell_{l,g,t}} p_{l,g}^A, & p_{l,g}^{B,+} &= e^{2\ell_{l,g,t}} p_{l,g}^B. \end{aligned} \quad (51)$$

The transformation law concerns freshly evaluated pressures; the EMA memories retain their temporal role in policy control. Positive ℓ reduces input-side relative pressure and increases output-side relative pressure. The $+0.25\pi_{l,g}$ term in equation 22 therefore raises $\tau_{l,g}$ when the input role has larger recent pressure, selecting the corrective sign in equation 51.

The curve-shape target follows from a separable role-pressure surrogate. After holding data- and orientation-dependent factors fixed within one group, write

$$\tilde{p}_{l,g}^A(\gamma) = \hat{d}_{l,g}^{\text{probe}} e^{-\gamma}, \quad \tilde{p}_{l,g}^B(\gamma) = \hat{o}_{l,g}^{\text{probe}} e^{\gamma}. \quad (52)$$

Their unique equality point is

$$\gamma_{l,g}^* = \frac{1}{2} \left(\log \hat{d}_{l,g}^{\text{probe}} - \log \hat{o}_{l,g}^{\text{probe}} \right), \quad (53)$$

which supplies the first term of $\tau_{l,g}$. This equality is a balance coordinate for the scalar proxy. Full data-dependent matrix gradients retain the orientation and data factors held fixed by the surrogate.

For nonzero matrices and unfloored RMS values,

$$\begin{aligned}\gamma_{l,g} &= \log \text{rms}(A_{l,g}) - \log \text{rms}(B_{l,g}), \\ \gamma_{l,g}^+ &= \gamma_{l,g} + 2\ell_{l,g,t}.\end{aligned}\tag{54}$$

Substituting the unclipped log step from equation 23 gives

$$\begin{aligned}\zeta_{l,g,t} &= s_{l,t}\chi_{l,g}^{\text{rs}}, \\ \gamma_{l,g}^+ &= (1 - \zeta_{l,g,t})\gamma_{l,g} + \zeta_{l,g,t}\tau_{l,g}.\end{aligned}\tag{55}$$

Since $\delta_l \in [0, 1]$, the evaluated depth factor lies in $[0.925, 1.075]$. Hence $0 \leq s_{l,t} \leq 0.5375$ and $\chi_{l,g}^{\text{rs}} \leq e^{0.125}$, so $0 \leq \zeta_{l,g,t} < 0.61$. The idealized unclipped coordinate is therefore a strict contraction toward $\tau_{l,g}$ whenever $\zeta_{l,g,t} > 0$. The 0.030 log-step clip bounds its magnitude and shortens moves that exceed the bound. With floored RMS values, $\gamma^+ = \gamma + 2\ell$ is approximate, and the implementation relies on the local bound in Appendix B.3.

The activity factor has the opposite sign in this speed control. Large $\alpha_{l,g}$ damps gradient-driven matrix motion in equation 18 and equation 20; here $\chi_{l,g}^{\text{rs}} = \exp(0.10\bar{\alpha}_{l,g})$ controls the speed of the post-step pair move, while the group and Muon controls determine gradient-driven matrix motion.