

RATIONAL MATRIX POLICY OPTIMIZATION FOR RATIONAL FEED-FORWARD LANGUAGE MODELS

Anonymous authors

Paper under double-blind review

ABSTRACT

1 INTRODUCTION

2 BACKGROUND AND RELATED WORK

3 RATIONAL MATRIX POLICY OPTIMIZATION

Let $\mathcal{L}(\theta)$ be the language-model training objective. For layer l , the RLB block projects the residual-stream input $x_l \in \mathbb{R}^d$ into G groups of width m ,

$$z_l = A_l x_l, \quad z_{l,g} = \text{group}_g(z_l), \quad (1)$$

normalizes each group by its RMS radius,

$$r_{l,g} = \sqrt{m^{-1} \|z_{l,g}\|_2^2 + \epsilon_{\text{rms}}}, \quad u_{l,g} = z_{l,g} / r_{l,g}, \quad (2)$$

and applies a learned rational map in normalized coordinates,

$$h_{l,g} = r_{l,g} R_{l,g}(u_{l,g}), \quad y_l = B_l \text{concat}_{g=1}^G h_{l,g}. \quad (3)$$

The input matrix A_l selects rational domains, the coefficients of $R_{l,g}$ shape the group response, and the output matrix B_l recombines rational features. MatrixPolicy uses this separation by partitioning parameters as

$$\theta = \theta_{\text{base}} \cup \theta_{\text{coeff}} \cup \{A_l, B_l\}_{l=1}^L. \quad (4)$$

The base parameters use the same optimizer as the non-rational baseline. Rational coefficients use a coefficient rule. MatrixPolicy is applied only to A_l and B_l .

In the homogeneous case $\epsilon_{\text{rms}} = 0$, and for nonzero group radii, each group has a positive scale gauge:

$$A_{l,g} \mapsto a_g A_{l,g}, \quad B_{l,g} \mapsto a_g^{-1} B_{l,g}, \quad a_g > 0. \quad (5)$$

The represented function is unchanged because $z_{l,g}$ and $r_{l,g}$ scale together, $u_{l,g}$ is unchanged, and the inverse scale in B_l cancels the scale in $h_{l,g}$. With a nonzero RMS floor this symmetry is only approximate, so we use the gauge as an optimizer-design principle rather than as a formal invariance claim. This matters because optimizer state is stored in coordinates, not in the quotient space of represented functions.

At step t , each matrix role $r \in \{\text{in}, \text{out}\}$ receives a role-depth factor $\rho_r(l)$ and an on-policy statistic factor $s_{\text{stat}}(l, r, t)$. The matrix update is

$$\begin{aligned} M_{l,r}^{t+1} = & M_{l,r}^t + \Delta_{\text{adapt}}(M_{l,r}; \eta_t a_{\text{mat}}(l, r, t)[1 - \mu(l, r, t)]) \\ & + \Delta_{\text{mat}}(M_{l,r}; \eta_t a_{\text{matnorm}} \mu(l, r, t)), \end{aligned} \quad (6)$$

where $M_{l,\text{in}} = A_l$, $M_{l,\text{out}} = B_l$, η_t is the shared base learning rate, Δ_{adapt} is an Adam-style adaptive step, and Δ_{mat} is a short early matrix-normalized step. The mixture coefficient $\mu(l, r, t)$ is

active only early in training. After the matrix update, MatrixPolicy computes a target log-ratio for each group and applies the bounded rebalance

$$A_{l,g} \leftarrow e^{\ell_g} A_{l,g}, \quad B_{l,g} \leftarrow e^{-\ell_g} B_{l,g}. \quad (7)$$

The update changes the representative used by the optimizer while preserving, or with $\epsilon_{\text{rms}} > 0$ approximately preserving, the represented block function. Full hyperparameter definitions are in Appendix A.

Algorithm 1: Rational Matrix Policy step. Given gradients $\nabla \mathcal{L}(\theta^t)$ and base learning rate η_t :

1. Update exponential moving averages of rational-group pressures.
2. Compute role-depth and group-stat multipliers for each RLB matrix.
3. Update non-rational parameters with the base optimizer and update rational coefficients with the coefficient rule.
4. Update each A_l and B_l with the MatrixPolicy mixture step.
5. Apply the bounded positive-gauge rebalance to each rational group.

4 EXPERIMENTS

5 DISCUSSION

6 CONCLUSION

REFERENCES

- Xiangning Chen, Chen Liang, Da Huang, Esteban Real, Kaiyuan Wang, Hieu Pham, Xuanyi Dong, Thang Lu, Cho-Jui Hsieh, Yifeng Lu, and Quoc V. Le. Symbolic Discovery of Optimization Algorithms. In *Advances in Neural Information Processing Systems*, 2023. URL https://papers.neurips.cc/paper_files/paper/2023/hash/9a39b4925e35cf447ccba8757137d84f-Abstract-Conference.html.
- Aaron Defazio, Xingyu Yang, Ahmed Khaled, Konstantin Mishchenko, Harsh Mehta, and Ashok Cutkosky. The Road Less Scheduled. In *Advances in Neural Information Processing Systems*, 2024. URL <https://neurips.cc/virtual/2024/poster/96925>.
- Diederik P. Kingma and Jimmy Ba. Adam: A Method for Stochastic Optimization. In *International Conference on Learning Representations*, 2015. URL <https://mlanthology.org/iclr/2015/kingma2015iclr-adam/>.
- Jeffrey Li, Alex Fang, Georgios Smyrnis, Maor Ivgi, Matt Jordan, Samir Gadre, Hritik Bansal, Etash Guha, et al. DataComp-LM: In Search of the Next Generation of Training Sets for Language Models. In *Advances in Neural Information Processing Systems, Datasets and Benchmarks Track*, 2024. doi: 10.52202/079017-0455. URL https://proceedings.neurips.cc/paper_files/paper/2024/hash/19e4ea30dded58259665db375885e412-Abstract-Datasets_and_Benchmarks_Track.html.
- Kaizhao Liang, Lizhang Chen, Bo Liu, and Qiang Liu. Cautious Optimizers: Improving Training with One Line of Code. In *International Conference on Learning Representations*, 2026. URL <https://openreview.net/forum?id=zBPZeRjfgu>.
- Hong Liu, Zhiyuan Li, David Hall, Percy Liang, and Tengyu Ma. Sophia: A Scalable Stochastic Second-order Optimizer for Language Model Pre-training. In *International Conference on Learning Representations*, 2024. URL https://proceedings.iclr.cc/paper_files/paper/2024/hash/06960915ba8674c7a898ec0b472b80ff-Abstract-Conference.html.

- Ilya Loshchilov and Frank Hutter. Decoupled Weight Decay Regularization. In *International Conference on Learning Representations*, 2019. URL <https://openreview.net/forum?id=Bkg6RiCqY7>.
- Yang Luo, Xiaozhe Ren, Zangwei Zheng, Zhuo Jiang, Xin Jiang, and Yang You. CAME: Confidence-guided Adaptive Memory Efficient Optimization. In *Proceedings of the 61st Annual Meeting of the Association for Computational Linguistics*, pp. 4442–4453, 2023. doi: 10.18653/v1/2023.acl-long.243. URL <https://aclanthology.org/2023.acl-long.243/>.
- Matteo Pagliardini, Pierre Ablin, and David Grangier. The AdEMAMix Optimizer: Better, Faster, Older. In *International Conference on Learning Representations*, 2025. URL <https://iclr.cc/virtual/2025/poster/28625>.
- Guilherme Penedo, Hynek Kydlicek, Loubna Ben Allal, Anton Lozhkov, Margaret Mitchell, Colin Raffel, Leandro von Werra, and Thomas Wolf. The FineWeb Datasets: Decanting the Web for the Finest Text Data at Scale. In *Advances in Neural Information Processing Systems, Datasets and Benchmarks Track*, 2024. doi: 10.52202/079017-0970. URL https://papers.nips.cc/paper/2024/hash/370df50ccfd8bde18f8f9c2d9151bda-Abstract-Datasets_and_Benchmarks_Track.html.
- Luca Soldaini, Rodney Kinney, Akshita Bhagia, Dustin Schwenk, David Atkinson, Russell Authur, Ben Bogin, Khyathi Chandu, et al. Dolma: An Open Corpus of Three Trillion Tokens for Language Model Pretraining Research. *arXiv preprint arXiv:2402.00159*, 2024. URL <https://arxiv.org/abs/2402.00159>.
- Nikhil Vyas, Depen Morwani, Rosie Zhao, Itai Shapira, David Brandfonbrener, Lucas Janson, and Sham Kakade. SOAP: Improving and Stabilizing Shampoo using Adam for Language Modeling. In *International Conference on Learning Representations*, 2025. URL https://proceedings.iclr.cc/paper_files/paper/2025/hash/e988664070e9591f93fdcf605f7dc623-Abstract-Conference.html.
- Kaiyue Wen, David Leo Wright Hall, Tengyu Ma, and Percy Liang. Fantastic Pretraining Optimizers and Where to Find Them. In *International Conference on Learning Representations*, 2026. URL <https://openreview.net/forum?id=2J51qUZ0iG>.
- Yushun Zhang, Congliang Chen, Ziniu Li, Tian Ding, Chenwei Wu, Diederik P. Kingma, Yinyu Ye, Zhi-Quan Luo, and Ruoyu Sun. Adam-mini: Use Fewer Learning Rates To Gain More. In *International Conference on Learning Representations*, 2025. URL https://proceedings.iclr.cc/paper_files/paper/2025/hash/45ae878717399e6f62d57c65f052cd46-Abstract-Conference.html.
- Jiawei Zhao, Zhenyu Zhang, Beidi Chen, Zhangyang Wang, Anima Anandkumar, and Yuandong Tian. GaLore: Memory-Efficient LLM Training by Gradient Low-Rank Projection. In *International Conference on Machine Learning*, 2024. URL <https://openreview.net/forum?id=hYHsrKDIX7>.

A MATRIXPOLICY DETAILS

Role-depth factors. For layers indexed by $l = 1, \dots, L$ with $L > 1$, define normalized layer depth $d_l = (l - 1)/(L - 1)$ and set

$$\rho_{\text{in}}(l) = \text{clip}(1 - 0.50(d_l - 0.5), 0.55, 1.40), \quad (8)$$

$$\rho_{\text{out}}(l) = \text{clip}(1 + 1.00(d_l - 0.5), 0.55, 1.40), \quad (9)$$

$$a_{\text{role}}(l, r) = \max(0.10, 1 + 1.20(\rho_r(l) - 1)). \quad (10)$$

Let $\xi_t = t/T$ be training progress and let $\sigma_{a,b}(\xi)$ denote the smoothstep gate from 0 to 1 over $[a, b]$. The statistic multiplier $a_{\text{stat}}(l, r, t) > 0$ is the layer/role aggregate of the group preconditioners defined below, normalized to one when group statistics are not active. The local adaptive scale is

$$a_{\text{mat}}(l, r, t) = \text{clip}(3.0 a_{\text{role}}(l, r) a_{\text{stat}}(l, r, t), 0.40, 4.0). \quad (11)$$

The matrix-normalized branch uses multiplier $a_{\text{matnorm}} > 0$ and mixture weight $\mu(l, r, t) \in [0, 1]$, a smoothstep-gated function of ξ_t , layer depth, matrix role, and $a_{\text{stat}}(l, r, t)$.

On-policy group statistics. All ratios and logarithms below are evaluated on strictly positive stabilized RMS, EMA, and gain quantities; the displayed formulas omit the corresponding epsilon floors. For group g , MatrixPolicy tracks relative gradient pressures

$$p_{\text{in},g} = \frac{\text{rms}(\nabla A_{l,g})}{\text{rms}(A_{l,g})}, \quad p_{\text{out},g} = \frac{\text{rms}(\nabla B_{l,g})}{\text{rms}(B_{l,g})}, \quad p_{\text{rat},g} = \text{rms}(\nabla \theta_{\text{coeff},g}). \quad (12)$$

With EMAs $q_{\text{in},g}$, $q_{\text{out},g}$, and $q_{\text{rat},g}$, we define

$$\pi_g = \log q_{\text{in},g} - \log q_{\text{out},g}, \quad \alpha_g = \log q_{\text{rat},g} - \frac{1}{2}(\log q_{\text{in},g} + \log q_{\text{out},g}). \quad (13)$$

Group-stat preconditioning. The group preconditioner is centered and clipped:

$$c_g = c_{\text{gain},g} c_{\text{pressure},g} c_{\text{activity},g}, \quad c_g \leftarrow c_g / \text{geomean}(c), \quad c_g \leftarrow \text{clip}(c_g, 0.75, 1.35). \quad (14)$$

Here $k_g > 0$ is the role-specific group gain used for matrix preconditioning: the derivative-RMS gain for A_l and the output-RMS gain for B_l . The factors are

$$c_{\text{gain},g} = \left(\frac{\text{geomean}(k)}{k_g} \right)^{0.20}, \quad c_{\text{pressure},g} = \exp(0.10 d_{\text{pressure},g}), \quad (15)$$

where $d_{\text{pressure},g} = -\pi_g$ for A_l and $+\pi_g$ for B_l . The activity damping term is

$$c_{\text{activity},g} = \exp[-0.20 \text{ReLU}((\alpha_g - 0.05)/0.45)]. \quad (16)$$

Gauge rebalance. Let

$$\gamma_g = \log \text{rms}(A_{l,g}) - \log \text{rms}(B_{l,g}). \quad (17)$$

The target log-ratio is

$$\tau_g = \frac{1}{2}(\log d_g^{\text{gain}} - \log o_g^{\text{gain}}) + \beta_{\text{pressure}}(\log q_{\text{in},g} - \log q_{\text{out},g}). \quad (18)$$

Here $d_g^{\text{gain}} > 0$ and $o_g^{\text{gain}} > 0$ are the derivative and output group-gain estimates, and β_{pressure} is the pressure balance coefficient. Let $s(t, l) \in [0, 1]$ be the rebalance schedule, $a_g^{\text{activity}} \in (0, 1]$ the activity damping multiplier, and $\ell_{\text{max}} > 0$ the maximum absolute log-scale move. The exact target step is $\ell_g^* = \frac{1}{2}(\tau_g - \gamma_g)$, while the applied bounded move is

$$\ell_g = \text{clip}(s(t, l) a_g^{\text{activity}} \ell_g^*, -\ell_{\text{max}}, \ell_{\text{max}}). \quad (19)$$

For true RMS values, the clipped update maps γ_g to $\gamma_g + 2\ell_g$ and is a bounded move toward τ_g rather than exact canonicalization unless clipping and schedule damping are inactive. With stabilized RMS values this identity is approximate; the implementation uses the bounded move as a controlled representative update for the RLB matrices.